

Programming OpenMP

Christian Terboven



Michael Klemm



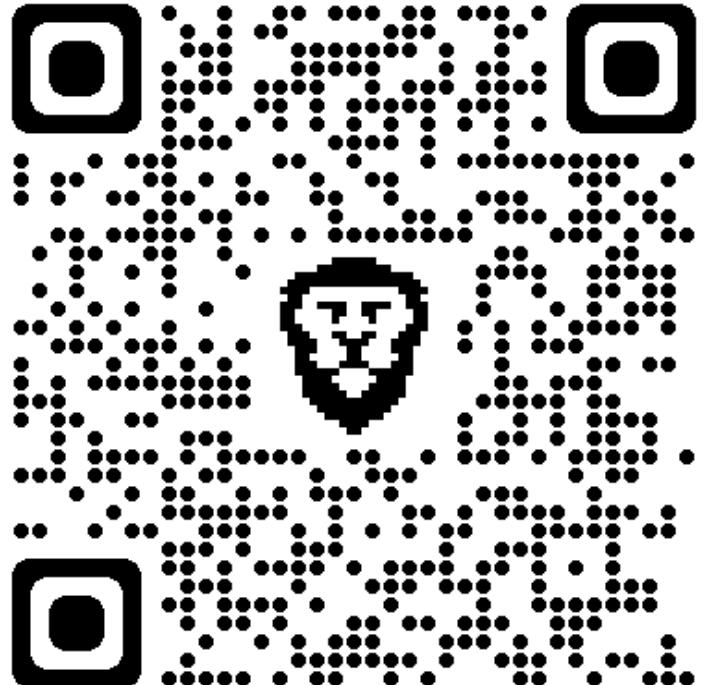
Agenda (tentative – tell us what else you need)

	Wed 07/12	Thu 08/12
08:50 – 09:00	Time to join the training	Time to join the training
09:00 – 10:30	Introduction	SIMD and GPUs Prog. / 1
10:30 – 10:45	Break	Break
10:45 – 12:15	Tasking / 1	GPU Prog. / 2
12:15 – 13:15	Break	Break
13:15 – 14:45	Tasking / 2	Misc. OpenMP Features
14:45 – 15:00	Break	Break
15:00 – 16:30	Hands-on / 1	Hands-on / 2

Lab: hands-on time

Material

- You can find all material on github.com:
 - Slide decks
 - Exercise tasks
 - Solutions
- <https://github.com/cterboven/OpenMP-tutorial-IT4I-2022>



Programming OpenMP

An Overview Of OpenMP

Christian Terboven

Michael Klemm



History

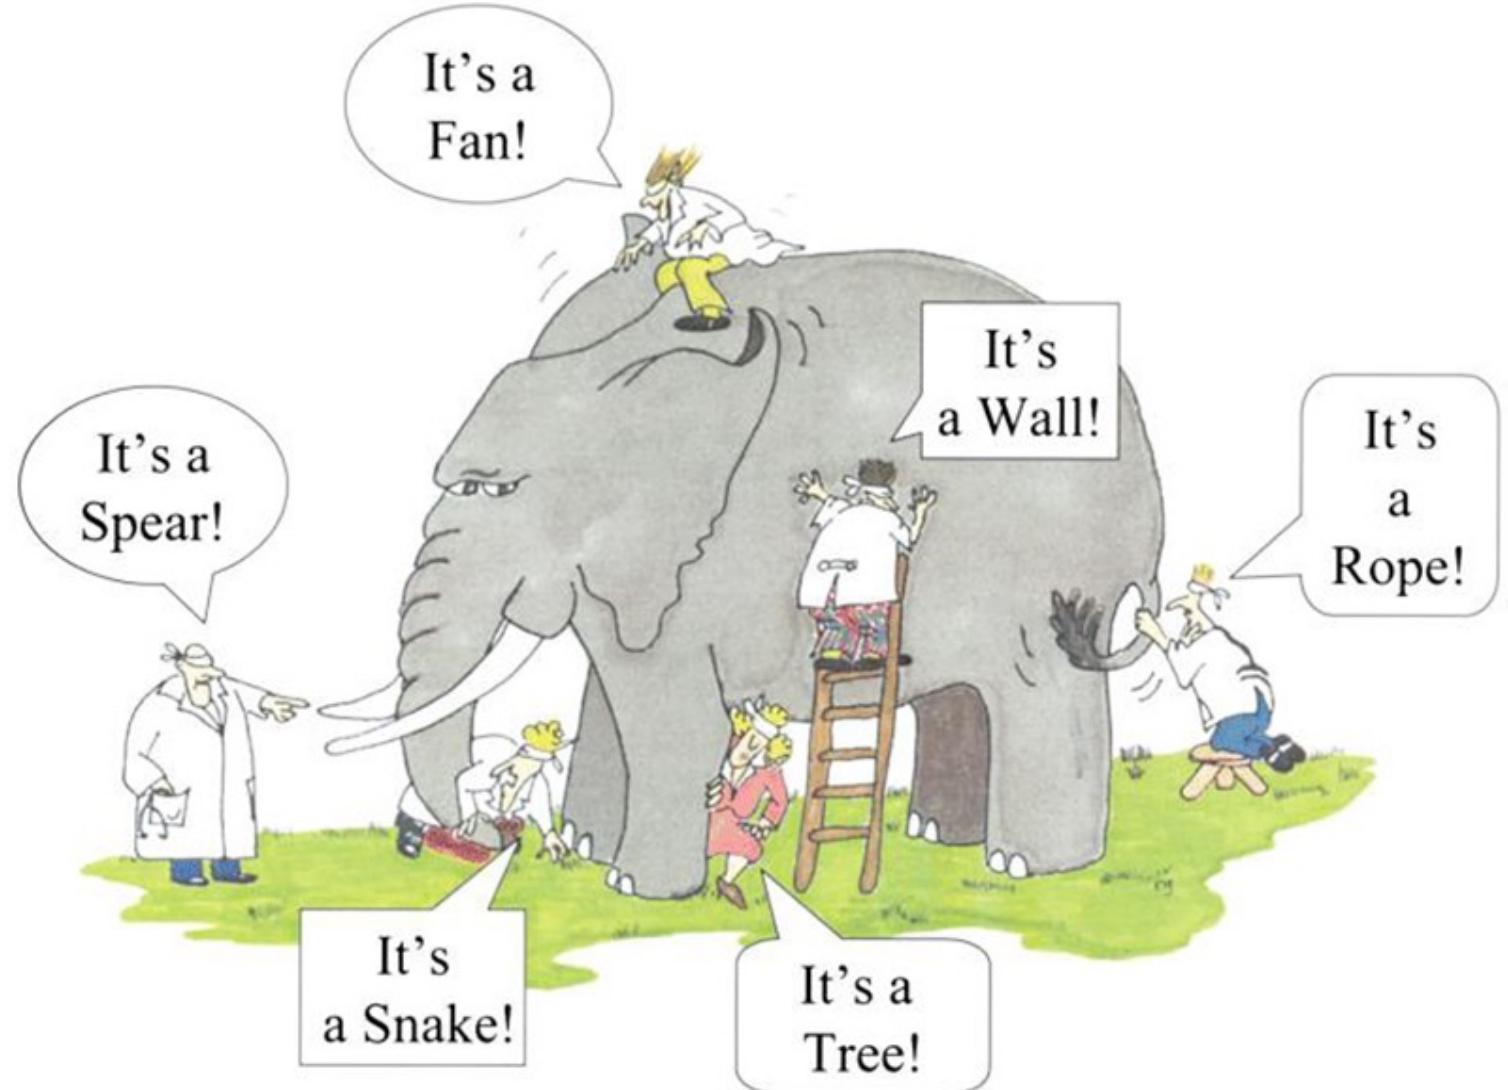
- De-facto standard for Shared-Memory Parallelization.
- 1997: OpenMP 1.0 for FORTRAN
- 1998: OpenMP 1.0 for C and C++
- 1999: OpenMP 1.1 for FORTRAN
- 2000: OpenMP 2.0 for FORTRAN
- 2002: OpenMP 2.0 for C and C++
- 2005: OpenMP 2.5 now includes both programming languages.
- 05/2008: OpenMP 3.0
- 07/2011: OpenMP 3.1
- 07/2013: OpenMP 4.0
- 11/2015: OpenMP 4.5
- 11/2018: OpenMP 5.0
- 11/2020: OpenMP 5.1
- 11/2021: OpenMP 5.2



<http://www.OpenMP.org>

What is OpenMP?

- Parallel Region & Worksharing
- Tasking
- SIMD / Vectorization
- Accelerator Programming
- ...



Get your C/C++ and Fortran Reference Guide!

Covers all of OpenMP 5.1/5.2!

OpenMP API 5.1

www.openmp.org

OpenMP 5.1 API Syntax Reference Guide

Page 1

Directives and Constructs

An OpenMP executable directive applies to the succeeding structured block. A structured-block is an OpenMP construct or a block of executable statements with a single entry at the top and a single exit at the bottom. OpenMP directives except `simd` and any declarative directive may not appear in Fortran PURE procedures.

Variant directives

metadirective [2.3.4][2.3.4]
A directive that can specify multiple directive variants, one of which may be automatically selected to replace the directive in the code to enable and to execute correctly.

```
!$omp omp metadirective [clause] [, clause] ...
!$omp omp begin metadirective [clause] [, clause] ...
!$omp end metadirective
!$omp metadirective [clause] [, clause] ...
!$omp begin metadirective [clause] [, clause] ...
!$omp end metadirective
!$omp clauses
when [control-selector-specification] [directive-variant]
declare variant [2.3.5][2.3.5]
```

declare variant declares a specialized variant of a base function and the context in which it is used.

```
!$omp omp declare variant [orient func-id] ...
clause [, clause] ...
!$omp omp begin variant [orient func-id] ...
clause [, clause] ...
!$omp end variant
!$omp declare variant [clause]
!$omp declare variant [func-name] [, ...
variant-proc-name] clause [, clause] ...

!$omp clauses
match [control-selector-specification]
ext_implementations-defined-requirement
assumption-clause
adjust-op: nothing, device, pr
append-op: nothing, interop-type [, interop-type] ...
!$omp variant-func-id: The name of a function variant that is available in C and C++, a template-id, or variant-proc-name: The name of a function variant that is a base language identifier.
```

dispatch [2.3.6]
Controls whether variant substitution occurs for a given call.

```
!$omp omp dispatch [clause] [, clause] ...
expression=stmt
!$omp dispatch [clause] [, clause] ...
stmt
!$omp clauses
depend [depend-modifier] , dependence-type : locator-lst
nowait
!$omp device [device-expression]
!$omp device [device-expression]
!$omp nocontrol [color-expression]
!$omp nocount [color-expression]
!$omp device [color-integer-expression]
!$omp nocontrol [color-integer-expression]
!$omp nocount [color-expression]
```

nothing [2.3.7]
Indicates explicitly that the intent is to have no effect.

```
!$omp omp nothing
!$omp clauses
```

OMPI120-01-OMP51

OpenMP API 5.1

www.openmp.org

Directives and Constructs (continued)

masked construct

masked [2.3.8][2.3.8]
Specifies a structured block that is executed by a subset of the threads of the current team. [In 5.0, this is the master construct, in which master replaces masked.]

```
!$omp masked [filter] [reorder-expression] ...
structured-block
!$omp masked ; filter[reorder-expression] ;
or
!$omp masked / filter[reorder-expression] /
strictly-structured-block
!$omp end masked
```

workshare

workshare [2.3.9][2.3.9]
Divides the execution of the enclosed structured block into separate units of work, each executed only once by iterations executed by each thread in the team.

```
!$omp workshare
!$omp workshare
!$omp workshare
!$omp workshare
!$omp end workshare [nowait]
```

Worksharing-loop construct

for do [2.3.10][2.3.10]
Specifies that the iterations of associated loops will be executed in parallel by threads in the team and the iterations executed by each thread can also be executed independently from the iterations of other threads.

```
!$omp for [clause] [, clause] ...
loop-body
!$omp do [clause] [, clause] ...
loop-body
!$omp end do [nowait]
```

scope construct

scope [2.3.11]
Defines a structured block that is executed by all threads in the team where additional OpenMP operations can be specified.

```
!$omp scope [clause] [, clause] ...
structured-block
!$omp scope [clause] [, clause] ...
loop-private [firstprivate] [lastprivate]
!$omp end scope [nowait]
!$omp scope [clause] [, clause] ...
collapse [n] [ordered [n]] [allocate [locator] [, locator] ...]
order [i : order-modifier] [concurrent]
reduction [reduction-modifier] [reduction-identifier] : [nowait]
```

parallel construct

parallel [2.3.12]
Creates a team of OpenMP threads that execute the region.

```
!$omp parallel [clause] [, clause] ...
structured-block
!$omp parallel [clause] [, clause] ...
loop-private [firstprivate] [lastprivate]
!$omp end parallel
!$omp parallel [clause] [, clause] ...
loop-private [firstprivate] [lastprivate]
!$omp end parallel
!$omp parallel [clause] [, clause] ...
loop-private [firstprivate] [lastprivate]
!$omp end parallel
```

distribute loop constructs

distribute [2.3.13][2.3.13]
Specifies loops which are executed by the initial teams.

```
!$omp omp distribute [clause] [, clause] ...
loop-body
!$omp distribute [clause] [, clause] ...
loop-body
!$omp end distribute
```

SIMD directives and constructs

simd [2.3.15][2.3.15]
Specifies that the associated structured block is executed by only one of the threads in the team.

```
!$omp omp simd [clause] [, clause] ...
structured-block
!$omp simd [clause] [, clause] ...
loop-body
!$omp end simd
```

parallel for and distribute parallel do

parallel for [2.3.16][2.3.16]
These constructs specify a loop that can be executed in parallel by multiple threads that are members of multiple teams.

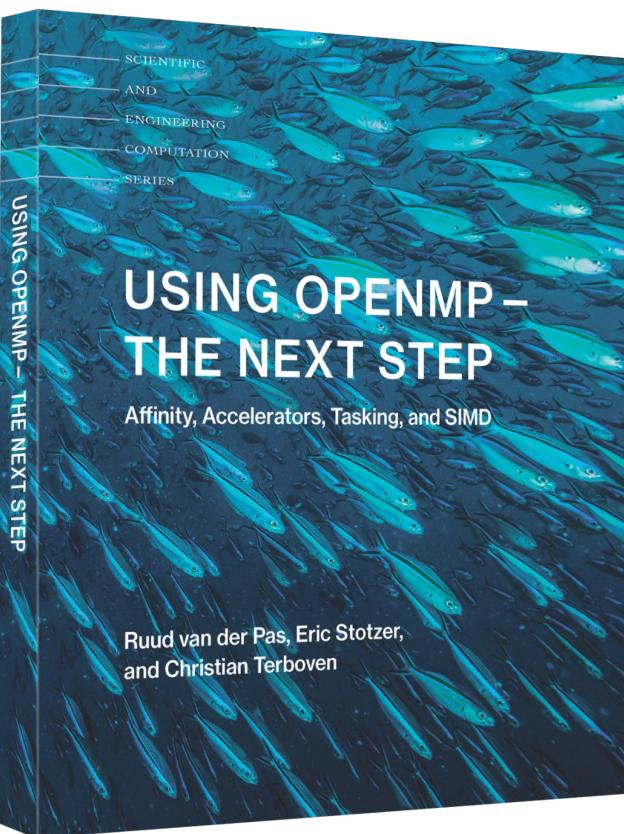
```
!$omp omp parallel for [clause] [, clause] ...
loop-body
!$omp distribute parallel do [clause] [, clause] ...
loop-body
!$omp end parallel do
```

OMPI120-01-OMP51

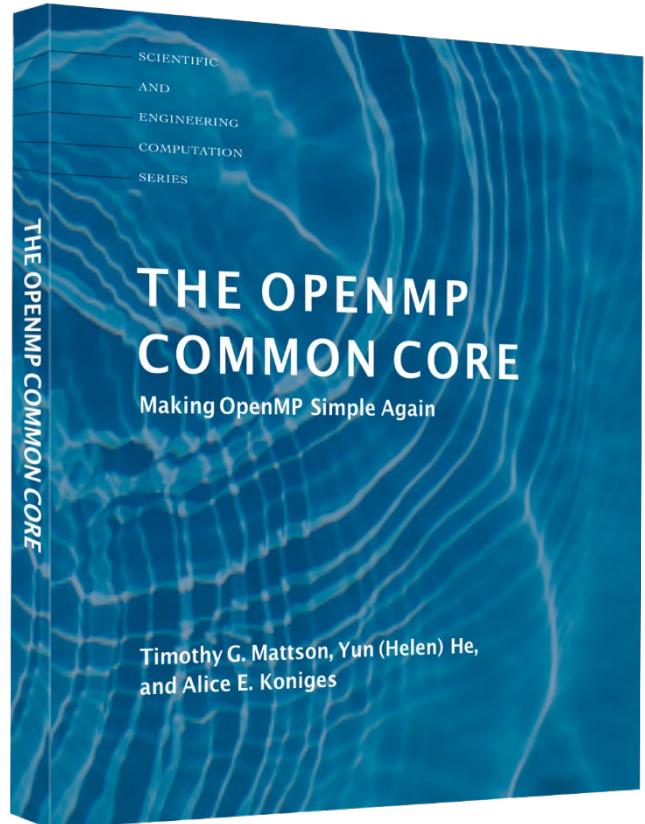
Recent Books About OpenMP



A printed copy of the 5.1 specifications, 2020



A book that covers all of the OpenMP 4.5 features, 2017



A new book about the OpenMP Common Core, 2019

Programming OpenMP

Parallel Region

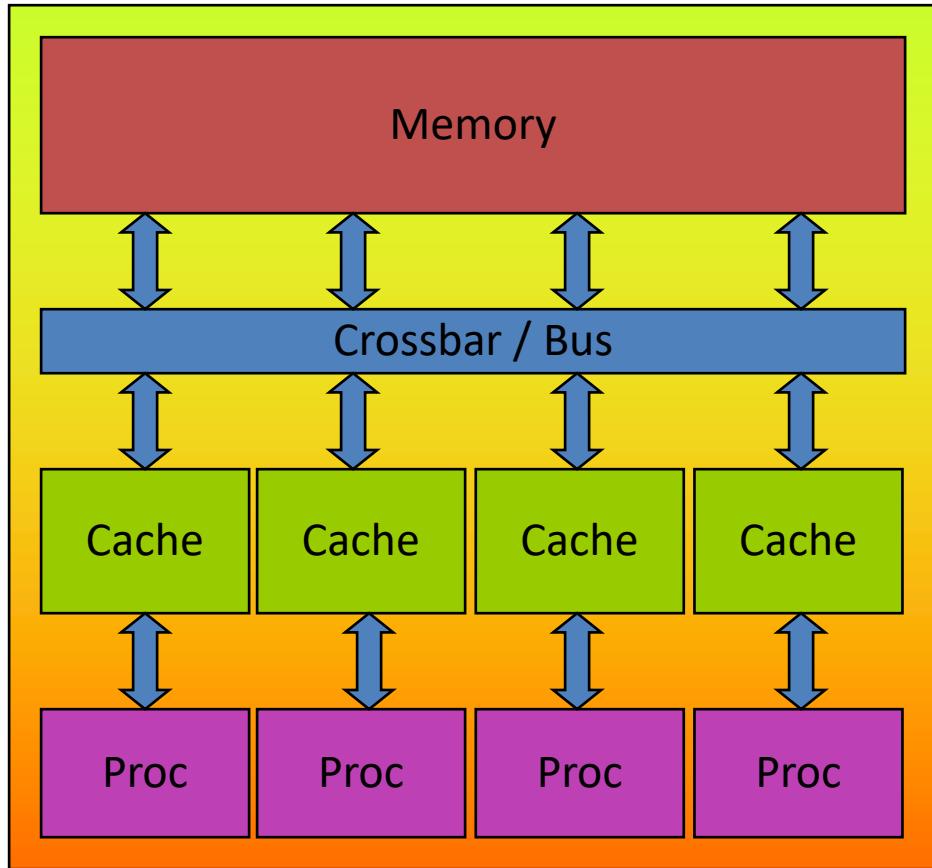
Christian Terboven

Michael Klemm



OpenMP's machine model

- OpenMP: Shared-Memory Parallel Programming Model.



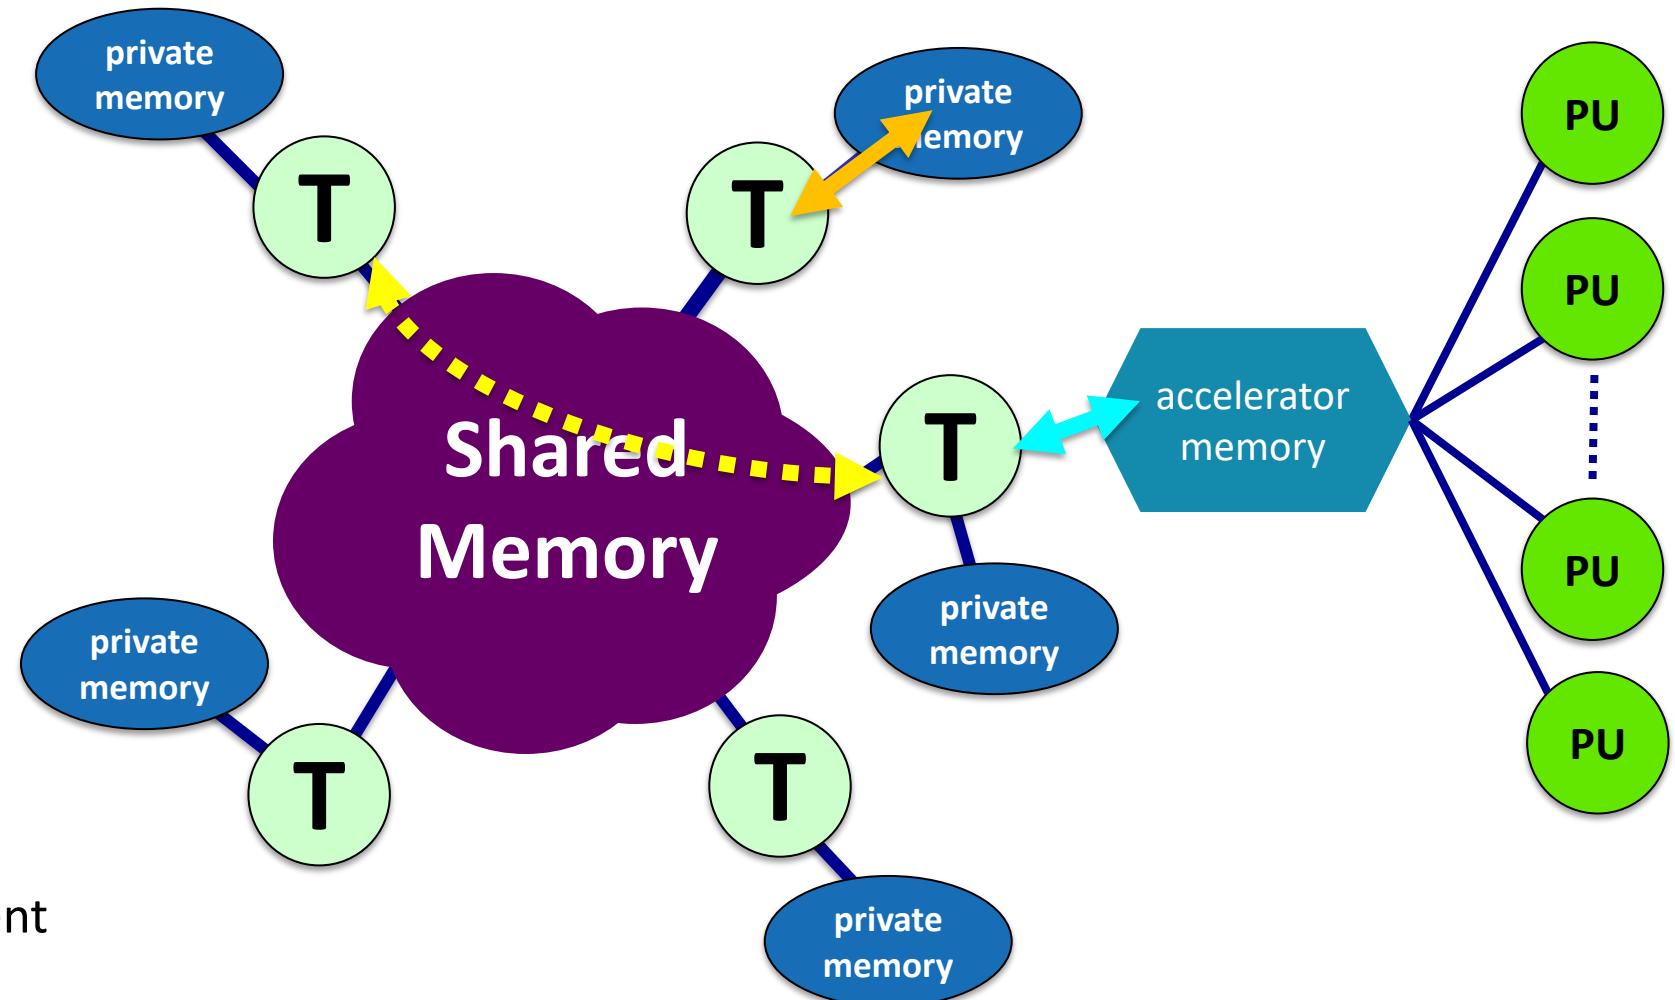
All processors/cores access a shared main memory.

Real architectures are more complex, as we will see later / as we

Parallelization in OpenMP employs multiple threads.

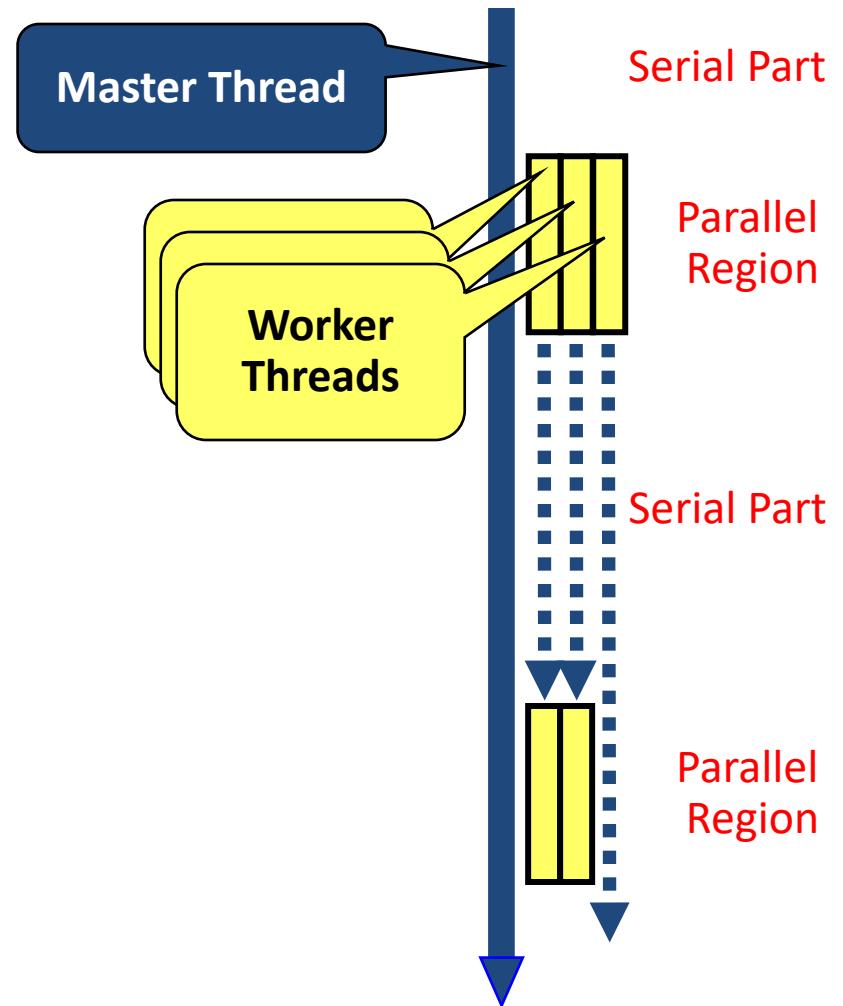
The OpenMP Memory Model

- All threads have access to the same, globally shared memory
- Data in private memory is only accessible by the thread owning this memory
- No other thread sees the change(s) in private memory
- Data transfer is through shared memory and is 100% transparent to the application



The OpenMP Execution Model

- OpenMP programs start with just one thread: The *Master*.
- *Worker* threads are spawned at *Parallel Regions*, together with the Master they form the *Team* of threads.
- In between Parallel Regions the Worker threads are put to sleep. The OpenMP *Runtime* takes care of all thread management work.
- Concept: *Fork-Join*.
- Allows for an incremental parallelization!



Parallel Region and Structured Blocks

- The parallelism has to be expressed explicitly.

C/C++

```
#pragma omp parallel
{
    ...
    structured block
    ...
}
```

Fortran

```
!$omp parallel
...
structured block
...
 !$omp end parallel
```

- Structured Block*

- Exactly one entry point at the top
- Exactly one exit point at the bottom
- Branching in or out is not allowed
- Terminating the program is allowed (abort / exit)

- Specification of number of threads:*

- Environment variable: OMP_NUM_THREADS=...
- Or: Via num_threads clause:
add num_threads (num) to the parallel construct

Starting OpenMP Programs on Linux

- From within a shell, global setting of the number of threads:

```
export OMP_NUM_THREADS=4  
./program
```

- From within a shell, one-time setting of the number of threads:

```
OMP_NUM_THREADS=4 ./program
```

Hello OpenMP World

Programming OpenMP

Worksharing

Christian Terboven

Michael Klemm



For Worksharing

- If only the *parallel* construct is used, each thread executes the Structured Block.
- Program Speedup: *Worksharing*
- OpenMP's most common Worksharing construct: *for*

C/C++

```
int i;  
#pragma omp for  
for (i = 0; i < 100; i++)  
{  
    a[i] = b[i] + c[i];  
}
```

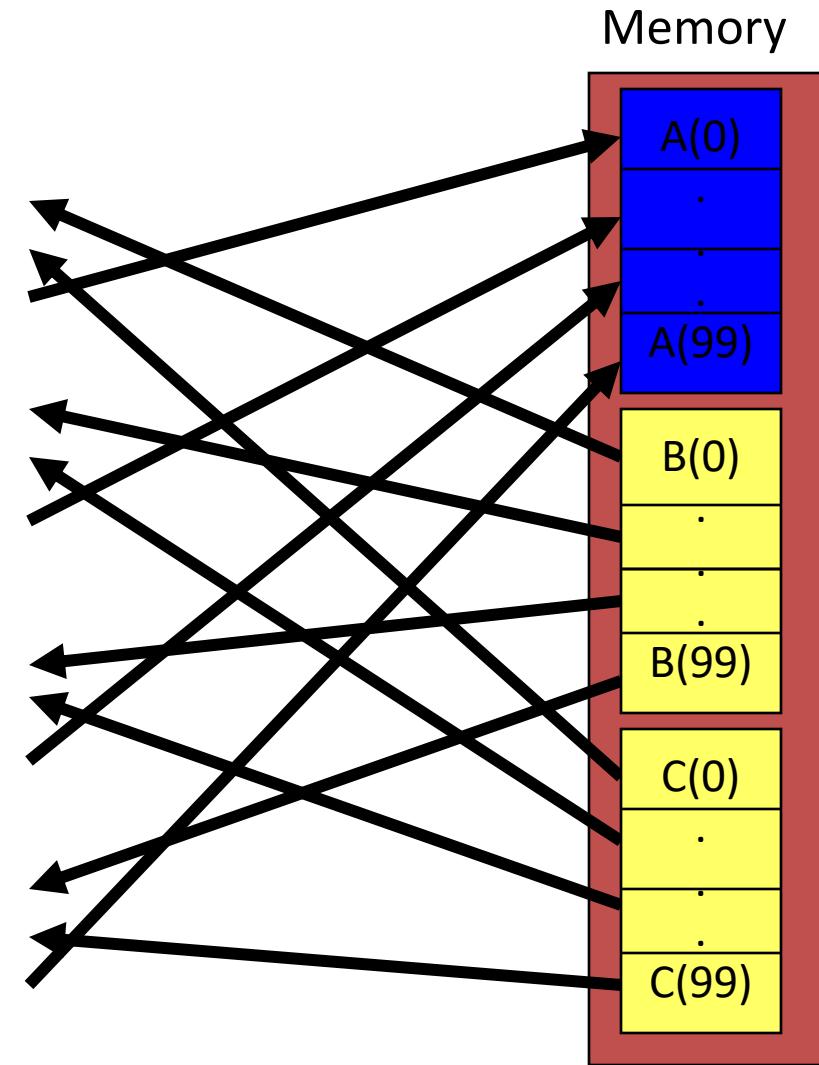
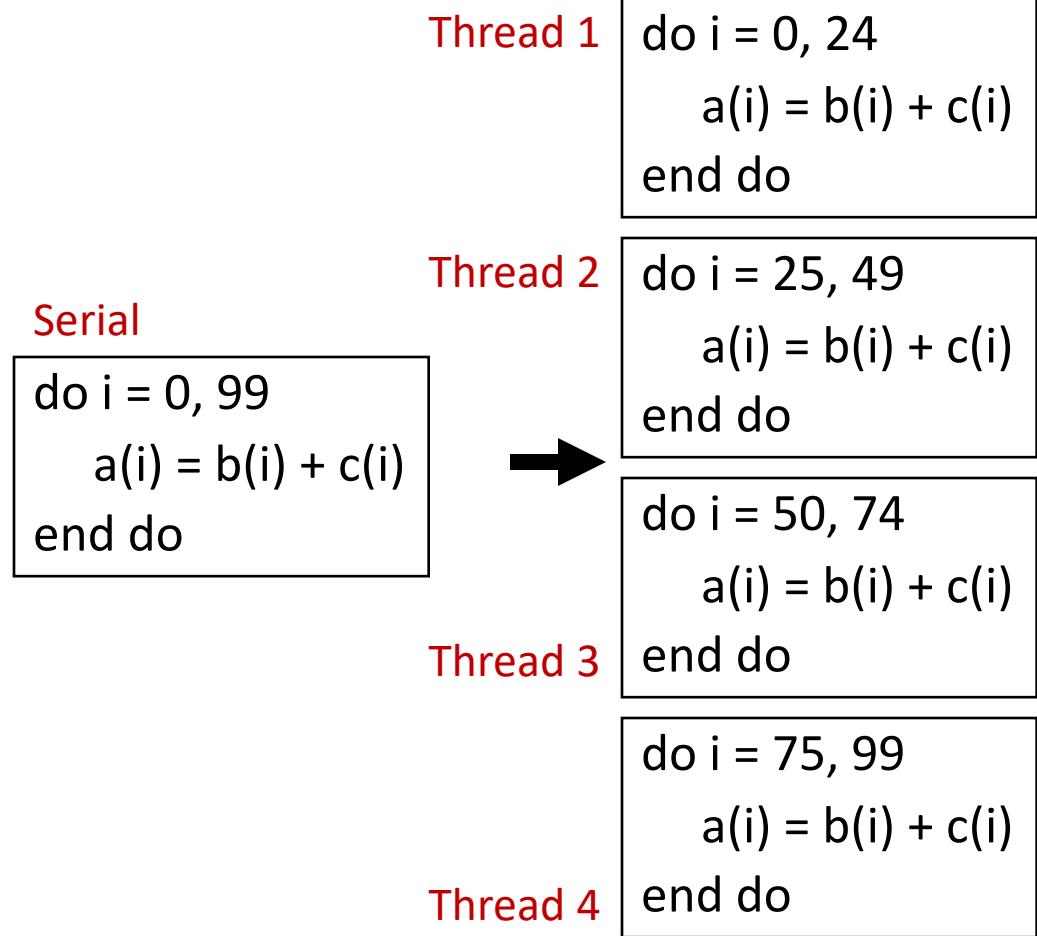
Fortran

```
INTEGER :: i  
!$omp do  
DO i = 0, 99  
    a[i] = b[i] + c[i]  
END DO
```

- Distribution of loop iterations over all threads in a Team.
 - Scheduling of the distribution can be influenced.
-
- Loops often account for most of a program's runtime!

Worksharing illustrated

Pseudo-Code
Here: 4 Threads



The Barrier Construct

- OpenMP barrier (implicit or explicit)
 - Threads wait until all threads of the current *Team* have reached the barrier

C/C++

```
#pragma omp barrier
```

- All worksharing constructs contain an implicit barrier at the end

The Single Construct

C/C++

```
#pragma omp single [clause]
... structured block ...
```

Fortran

```
!$omp single [clause]
... structured block ...
!$omp end single
```

- The **single** construct specifies that the enclosed structured block is executed by only one thread of the team.
 - It is up to the runtime which thread that is.
- Useful for:
 - I/O
 - Memory allocation and deallocation, etc. (in general: setup work)
 - Implementation of the single-creator parallel-executor pattern as we will see later...

The Master Construct is going to be removed with OpenMP 6.0 (2025)

C/C++

```
#pragma omp master[clause]
... structured block ...
```

Fortran

```
!$omp master[clause]
... structured block ...
!$omp end master
```

- The `master` construct specified that the enclosed structured block is executed only by the master thread of a team.
 - Note: The master construct was no worksharing construct and does not contain an implicit barrier at the end.
- Replacement: see the masked construct later on.

Vector Addition

Influencing the For Loop Scheduling / 1

- *for*-construct: OpenMP allows to influence how the iterations are scheduled among the threads of the team, via the *schedule* clause:
 - `schedule(static [, chunk])`: Iteration space divided into blocks of chunk size, blocks are assigned to threads in a round-robin fashion. If chunk is not specified: #threads blocks.
 - `schedule(dynamic [, chunk])`: Iteration space divided into blocks of chunk (not specified: 1) size, blocks are scheduled to threads in the order in which threads finish previous blocks.
 - `schedule(guided [, chunk])`: Similar to dynamic, but block size starts with implementation-defined value, then is decreased exponentially down to chunk.
- Default is `schedule(static)`.

■ Static Schedule

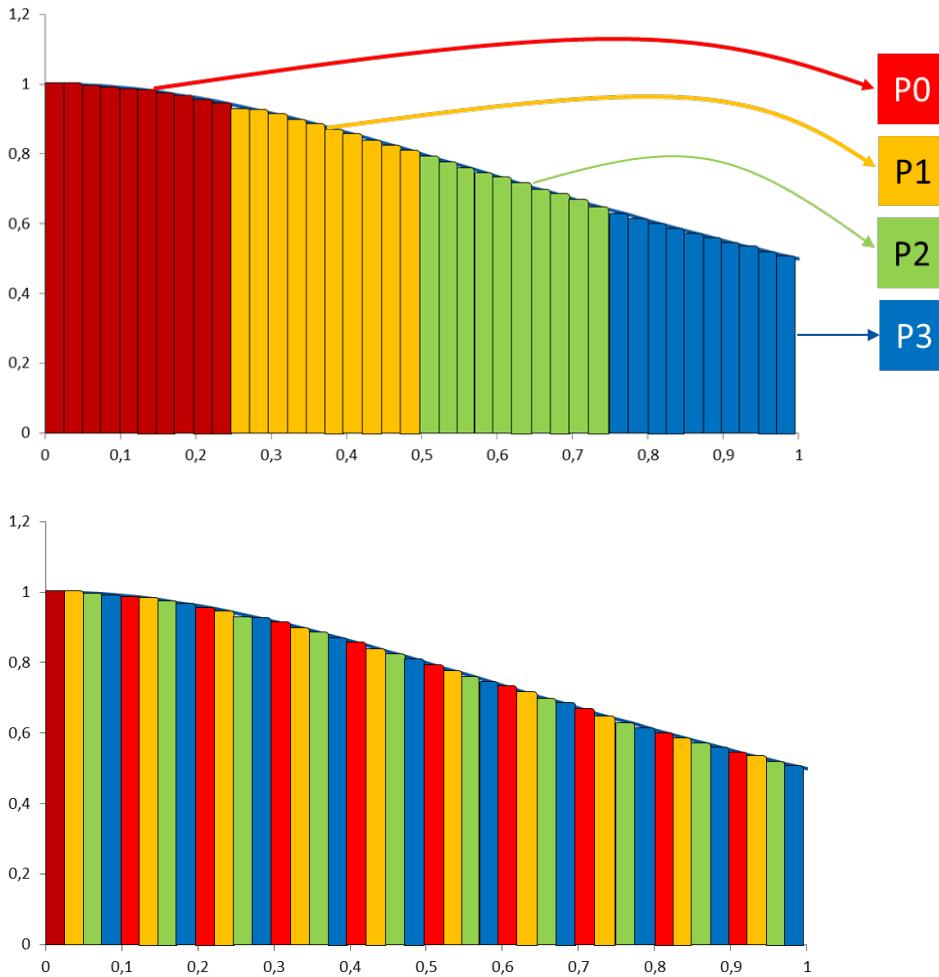
- `schedule(static [, chunk])`
- Decomposition depending on chunksize
- Equal parts of size ‘chunksize’ distributed in round-robin fashion

■ Pros?

- No/low runtime overhead

■ Cons?

- No dynamic workload balancing



Influencing the For Loop Scheduling / 3

- Dynamic schedule
 - `schedule(dynamic [, chunk])`
 - Iteration space divided into blocks of chunk size
 - Threads request a new block after finishing the previous one
 - Default chunk size is 1
- Pros ?
 - Workload distribution
- Cons?
 - Runtime Overhead
 - Chunk size essential for performance
 - No NUMA optimizations possible

Synchronization Overview

- Can all loops be parallelized with `for`-constructs? No!
 - Simple test: If the results differ when the code is executed backwards, the loop iterations are not independent.
BUT: This test alone is not sufficient:

C/C++

```
int i, int s = 0;  
  
#pragma omp parallel for  
for (i = 0; i < 100; i++)  
{  
    s = s + a[i];  
}
```

- *Data Race*: If between two synchronization points at least one thread writes to a memory location from which at least one other thread reads, the result is not deterministic (race condition).

Synchronization: Critical Region

- A *Critical Region* is executed by all threads, but by only one thread simultaneously (*Mutual Exclusion*).

C/C++

```
#pragma omp critical (name)
{
    ... structured block ...
}
```

- Do you think this solution scales well?

C/C++

```
int i, s = 0;
#pragma omp parallel for
for (i = 0; i < 100; i++)
{
    #pragma omp critical
    {
        s = s + a[i];
    }
}
```

Programming OpenMP

Scoping

Christian Terboven

Michael Klemm

RWTH AACHEN
UNIVERSITY
OpenMP®

Scoping Rules

- Managing the Data Environment is the challenge of OpenMP.
- *Scoping* in OpenMP: Dividing variables in *shared* and *private*:
 - *private*-list and *shared*-list on Parallel Region
 - *private*-list and *shared*-list on Worksharing constructs
 - General default is *shared* for Parallel Region, *firstprivate* for Tasks.
Tasks are introduced later
 - Loop control variables on *for*-constructs are *private*
 - Non-static variables local to Parallel Regions are *private*
 - *private*: A new uninitialized instance is created for the task or each thread executing the construct
 - *firstprivate*: Initialization with the value before encountering the construct
 - *lastprivate*: Value of last loop iteration is written back to Master
 - Static variables are *shared*

Privatization of Global/Static Variables

- Global / static variables can be privatized with the *threadprivate* directive
 - One instance is created for each thread
 - Before the first parallel region is encountered
 - Instance exists until the program ends
 - Does not work (well) with nested Parallel Region
 - Based on thread-local storage (TLS)
 - TlsAlloc (Win32-Threads), pthread_key_create (Posix-Threads), keyword `__thread` (GNU extension)

C/C++

```
static int i;  
#pragma omp threadprivate(i)
```

Fortran

```
SAVE INTEGER :: i  
!$omp threadprivate(i)
```

Privatization of Global/Static Variables

- Global / static variables can be privatized with the *threadprivate* directive
 - One instance is created for each thread
 - Before the first parallel region is encountered
 - Instance exists until the program ends
 - Does not work (well) with nested Parallel Region
 - Based on thread-local storage (TLS)
 - TlsAlloc (Win32-Threads), pthread_key_create (Posix Threads), keyword __thread (GNU extension)

C/C++

```
static int i;  
#pragma omp threadprivate(i)
```

Fortran

```
SAVE INTEGER :: i  
!$omp threadprivate(i)
```

Really: try to avoid the use of *threadprivate* and static variables!

Back to our example

C/C++

```
int i, s = 0;  
#pragma omp parallel for  
for (i = 0; i < 100; i++)  
{  
    #pragma omp critical  
        { s = s + a[i]; }  
}
```

It's your turn: Make It Scale!

```
#pragma omp parallel
```

```
{
```

```
#pragma omp for
```

```
for (i = 0; i < 99; i++)  
{
```

```
    s = s + a[i];
```

```
}
```

```
} // end parallel
```

```
do i = 0, 24  
    s = s + a(i)  
end do
```

```
do i = 0, 99  
    s = s + a(i)  
end do
```



```
do i = 25, 49  
    s = s + a(i)  
end do
```

```
do i = 50, 74  
    s = s + a(i)  
end do
```

```
do i = 75, 99  
    s = s + a(i)  
end do
```

(done)

```
#pragma omp parallel
{
    double ps = 0.0;      // private variable
#pragma omp for
for (i = 0; i < 99; i++)
{
    ps = ps + a[i];
}
#pragma omp critical
{
    s += ps;
}
} // end parallel
```

do i = 0, 99
s = s + a(i)
end do

do i = 0, 24
s₁ = s₁ + a(i)
end do
s = s + s₁

do i = 25, 49
s₂ = s₂ + a(i)
end do
s = s + s₂

do i = 50, 74
s₃ = s₃ + a(i)
end do
s = s + s₃

do i = 75, 99
s₄ = s₄ + a(i)
end do
s = s + s₄

The Reduction Clause

- In a *reduction*-operation the operator is applied to all variables in the list. The variables have to be *shared*.
 - reduction(operator:list)
 - The result is provided in the associated reduction variable

C/C++

```
int i, s = 0;  
  
#pragma omp parallel for reduction(+:s)  
for(i = 0; i < 99; i++)  
{  
    s = s + a[i];  
}
```

- Possible reduction operators with initialization value:
 - + (0), * (1), - (0), & (~0), | (0), && (1), || (0), ^ (0), min (largest number), max (least number)
- Remark: OpenMP also supports user-defined reductions (not covered here)

PI

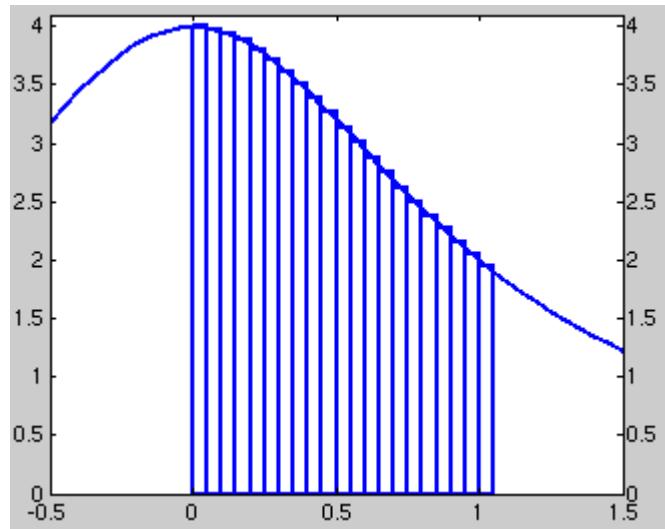
Example: Pi (1/2)

```
double f(double x)
{
    return (4.0 / (1.0 + x*x));
}

double CalcPi (int n)
{
    const double fH = 1.0 / (double) n;
    double fSum = 0.0;
    double fX;
    int i;

#pragma omp parallel for
    for (i = 0; i < n; i++)
    {
        fX = fH * ((double)i + 0.5);
        fSum += f(fX);
    }
    return fH * fSum;
}
```

$$\pi = \int_0^1 \frac{4}{1+x^2}$$



Example: Pi (2/2)

```

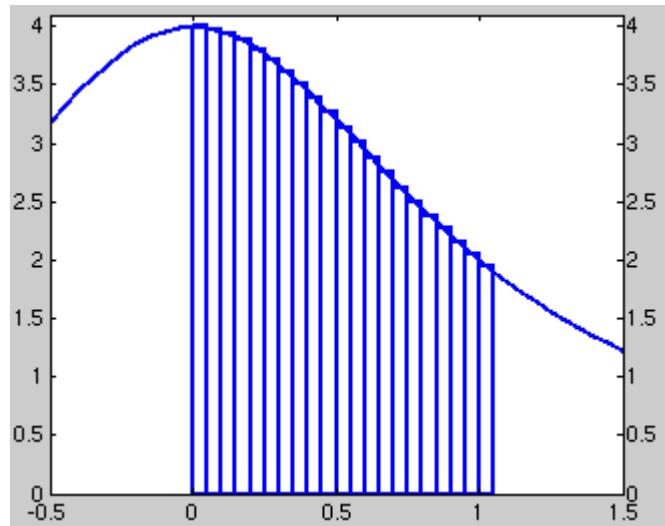
double f(double x)
{
    return (4.0 / (1.0 + x*x));
}

double CalcPi (int n)
{
    const double fH = 1.0 / (double) n;
    double fSum = 0.0;
    double fX;
    int i;

#pragma omp parallel for private(fX,i) reduction(+:fSum)
    for (i = 0; i < n; i++)
    {
        fX = fH * ((double)i + 0.5);
        fSum += f(fX);
    }
    return fH * fSum;
}

```

$$\pi = \int_0^1 \frac{4}{1+x^2}$$



PI

Programming OpenMP

Tasking

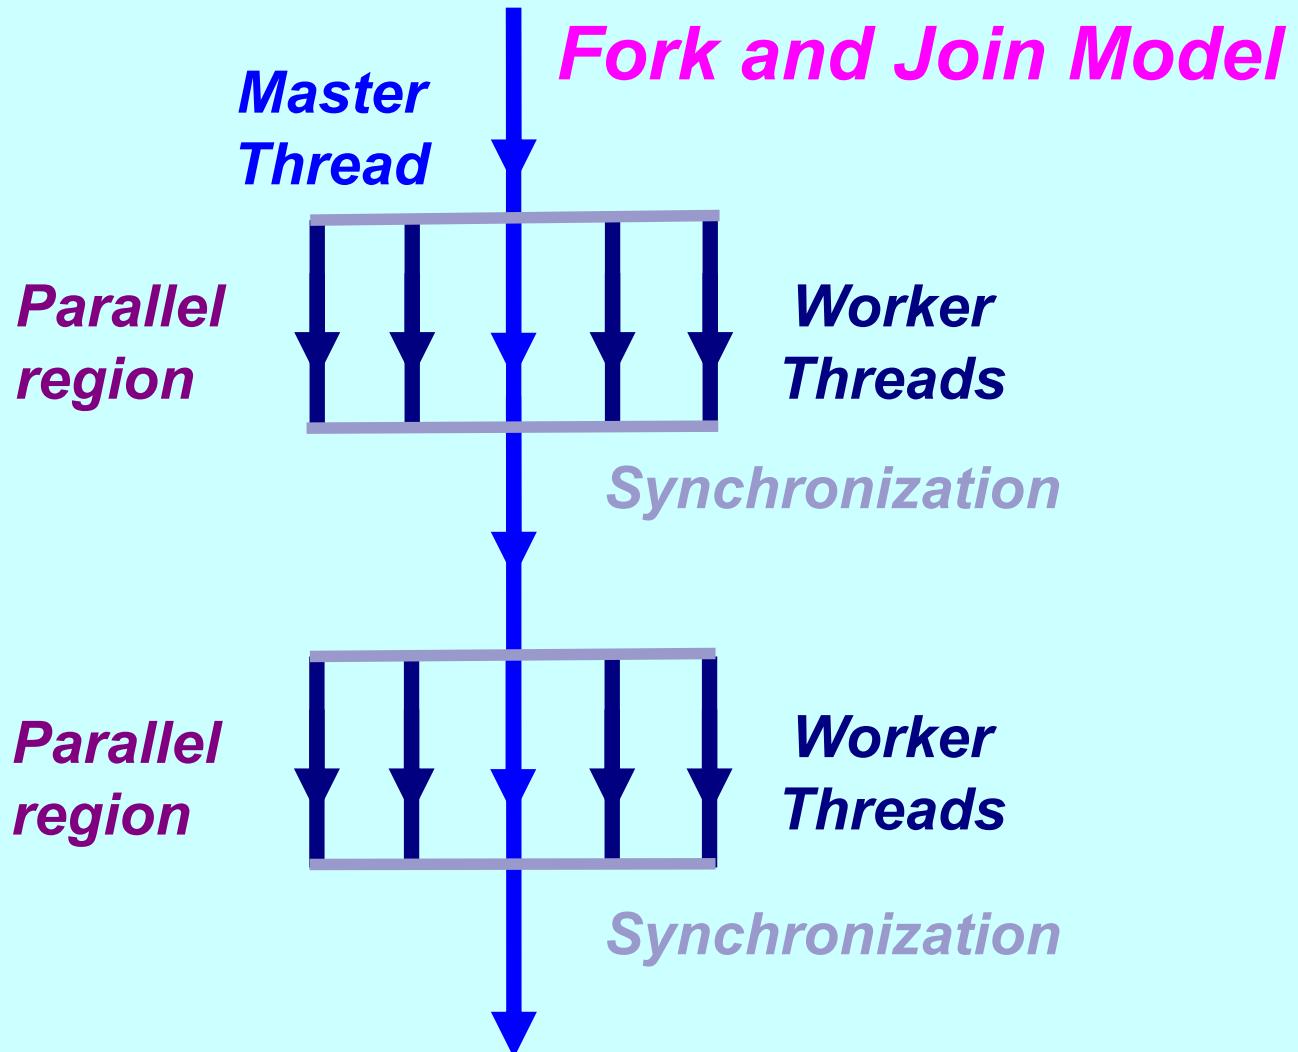
Christian Terboven

Michael Klemm

**RWTHAACHEN
UNIVERSITY**
OpenMP[®]

OpenMP Review

The OpenMP Execution Model



```
#pragma omp parallel  
{  
    ....  
}
```

```
#pragma omp parallel  
{  
    ....  
}
```

The Worksharing Constructs

- ***The work is distributed over the threads***
- ***Must be enclosed in a parallel region***
- ***Must be encountered by all threads in the team, or none at all***
- ***No implied barrier on entry***
- ***Implied barrier on exit (unless the nowait clause is specified)***
- ***A work-sharing construct does not launch any new threads***

```
#pragma omp for  
{  
    ....  
}
```

```
#pragma omp sections  
{  
    ....  
}
```

```
#pragma omp single  
{  
    ....  
}
```

The Single and Master Directives

- Single: only one thread in the team executes the code enclosed

```
#pragma omp single [private] [firstprivate] \
    [copyprivate] [nowait]
{
    <code-block>
}
```

- Master: the master thread executes the code enclosed

```
#pragma omp master
{<code-block>}
```

*There is no implied
barrier on entry or
exit !*



The OpenMP Barrier

- Several constructs have an implied barrier
 - This is another safety net (has implied flush by the way)
 - the “nowait” clause
- This can help fine tuning the application
 - But you'd better know what you're doing
- The explicit barrier comes in quite handy then

```
#pragma omp barrier
```



Tasking Motivation

Sudoku for Lazy Computer Scientists

- Lets solve Sudoku puzzles with brute multi-core force

	6					8	11			15	14				16
15	11			16	14			12			6				
13		9	12			3	16	14		15	11	10			
2	16		11	15	10	1									
	15	11	10		16	2	13	8	9	12					
12	13			4	1	5	6	2	3				11	10	
5		6	1	12		9	15	11	10	7	16				3
	2			10		11	6		5		13				9
10	7	15	11	16			12	13							6
9					1		2		16	10					11
1	4	6	9	13		7		11		3	16				
16	14		7	10	15	4	6	1			13	8			
11	10		15			16	9	12	13		1	5	4		
	12		1	4	6	16			11	10					
	5		8	12	13	10			11	2					14
3	16		10		7		6				12				

- (1) Search an empty field
- (2) Try all numbers:
 - (2 a) Check Sudoku
 - If invalid: skip
 - If valid: Go to next field
- Wait for completion



Parallel Brute-force Sudoku

- This parallel algorithm finds all valid solutions

	6					8	11			15	14			16
15	11			16	14			12			6			
13		9	12			3	16	14		15	11	10		
2	16		11	15	10	1								
	15	11	10		16	2	13	8	9	12				
12	13			4	1	5	6	2	3				11	10
5		6	1	12		9	15	11	10	7	16			3
	2			10		11	6		5		13			9
10	7	15	11	16			12	13						6
9					1		2		16	10				11
1	4	6	9	13		7	11		3	16				
16	14		7	10	15	4	6	1			13	8		
11	10		15			16	9	12	13		1	5	4	
	12		1	4	6	16			11	10				
	5		8	12	13	10		11	2				14	
3	16		10		7		6				12			

- (1) Search an empty field

first call contained in a
`#pragma omp parallel`
`#pragma omp single`
such that one tasks starts the
execution of the algorithm

- (2) Try all numbers:

- (2 a) Check Sudoku
- If invalid: skip
- If valid: Go to next field

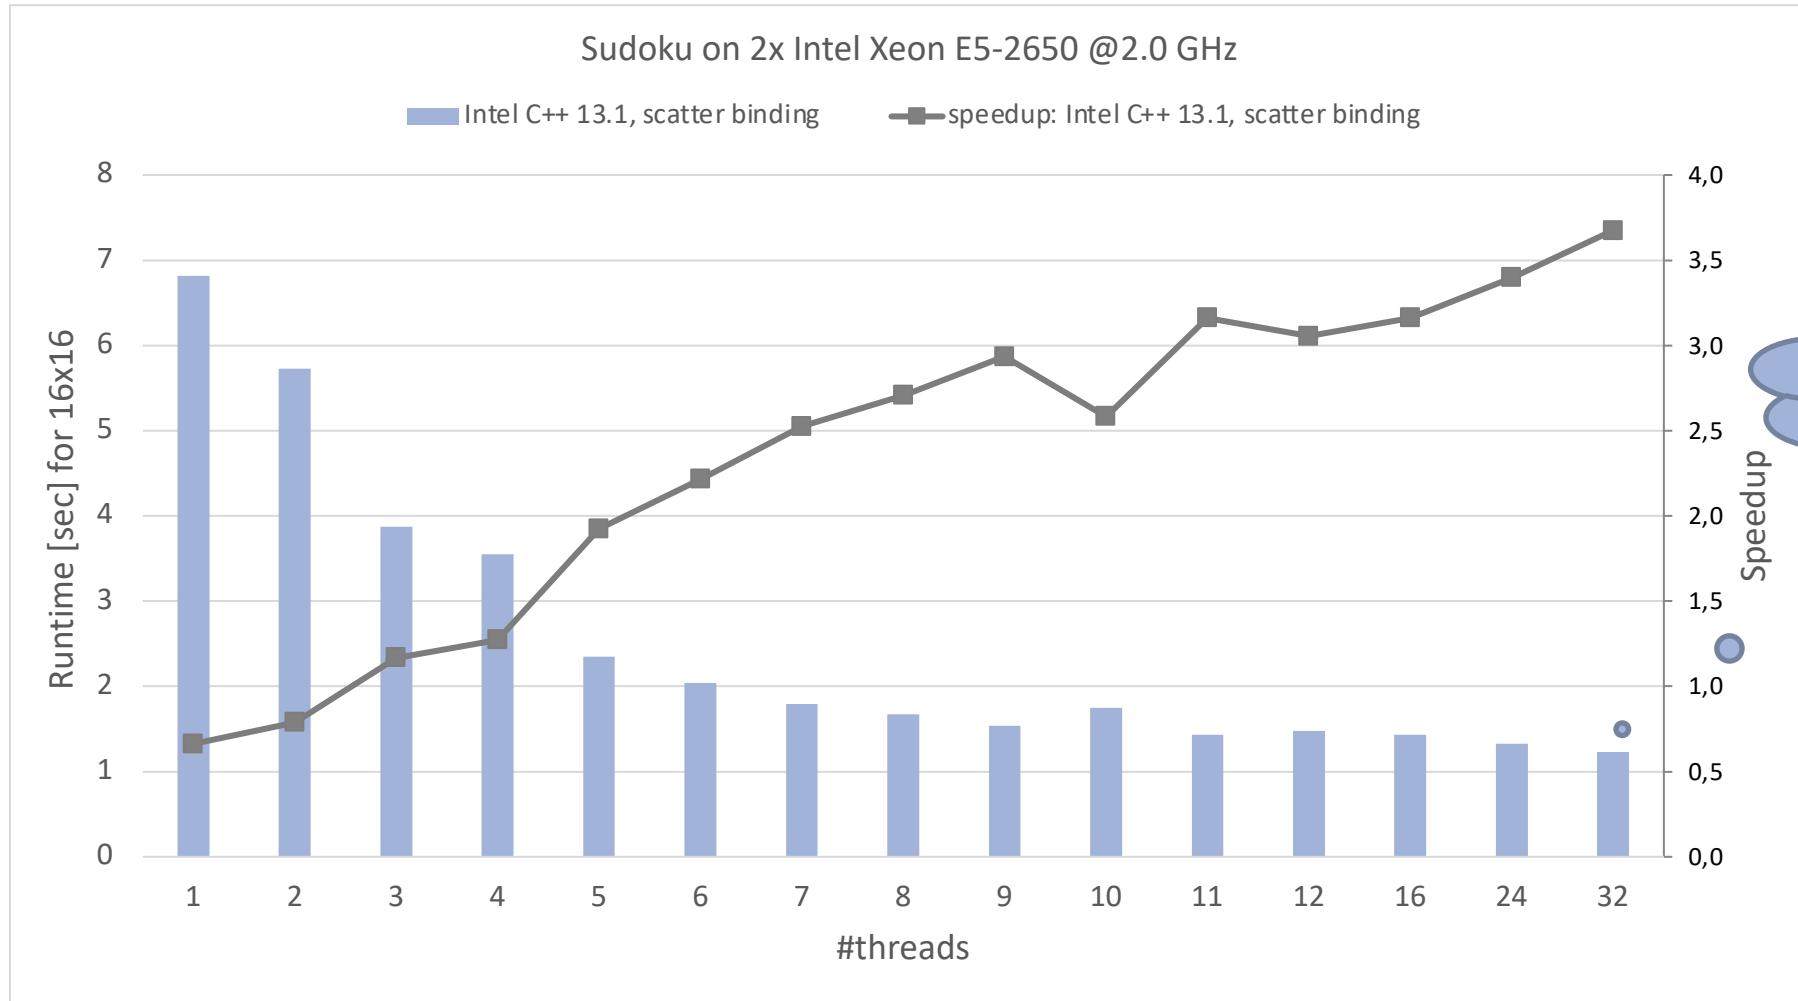
`#pragma omp task`
needs to work on a new copy
of the Sudoku board

- Wait for completion

`#pragma omp taskwait`
wait for all child tasks



Performance Evaluation



Is this the best
we can do?

Tasking Overview



What is a task in OpenMP?

- Tasks are work units whose execution
 - may be deferred or...
 - ... can be executed immediately
- Tasks are composed of
 - **code** to execute, a **data** environment (initialized at creation time), internal **control** variables (ICVs)
- Tasks are created...
 - ... when reaching a parallel region → implicit tasks are created (per thread)
 - ... when encountering a task construct → explicit task is created
 - ... when encountering a taskloop construct → explicit tasks per chunk are created
 - ... when encountering a target construct → target task is created

Tasking execution model

- Supports unstructured parallelism

→ unbounded loops

```
while ( <expr> ) {
    ...
}
```

→ recursive functions

```
void myfunc( <args> )
{
    ...; myfunc( <newargs> ); ...
}
```

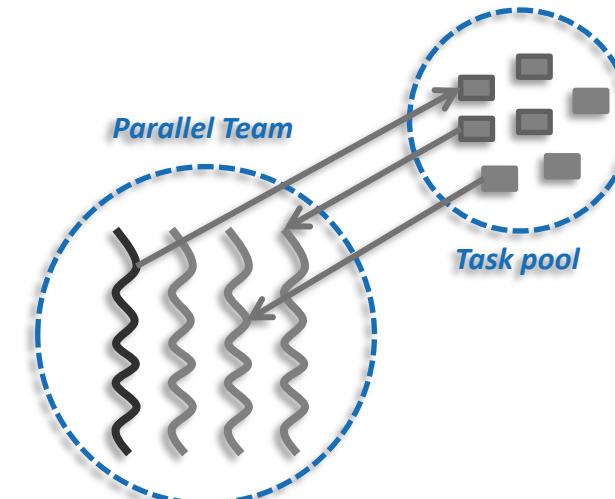
- Several scenarios are possible:

→ single creator, multiple creators, nested tasks (tasks & WS)

- All threads in the team are candidates to execute tasks

- Example (unstructured parallelism)

```
#pragma omp parallel
#pragma omp master
while (elem != NULL) {
    #pragma omp task
    compute(elem);
    elem = elem->next;
}
```



The task construct

- Deferring (or not) a unit of work (executable for any member of the team)

```
#pragma omp task [clause[,] clause]...
{structured-block}
```

```
!$omp task [clause[,] clause]...
...structured-block...
 !$omp end task
```

- Where clause is one of:

- private(list)
- firstprivate(list)
- shared(list)
- default(shared | none)
- in_reduction(r-id: list)

Data Environment

- allocate([allocator:] list)
- detach(event-handler)

Miscellaneous

- if(scalar-expression)
- mergeable
- final(scalar-expression)

Cutoff Strategies

- depend(dep-type: list)

Synchronization

- untied
- priority(priority-value)
- affinity(list)

Task Scheduling



Task scheduling: tied vs untied tasks

- Tasks are tied by default (when no untied clause present)
 - tied tasks are executed always by the same thread (not necessarily creator)
 - tied tasks may run into performance problems
- Programmers may specify tasks to be untied (relax scheduling)

```
#pragma omp task untied  
{structured-block}
```

- can potentially switch to any thread (of the team)
- bad mix with thread based features: thread-id, threadprivate, critical regions...
- gives the runtime more flexibility to schedule tasks
- but most of OpenMP implementations doesn't "honor" untied ☹



Task scheduling: taskyield directive

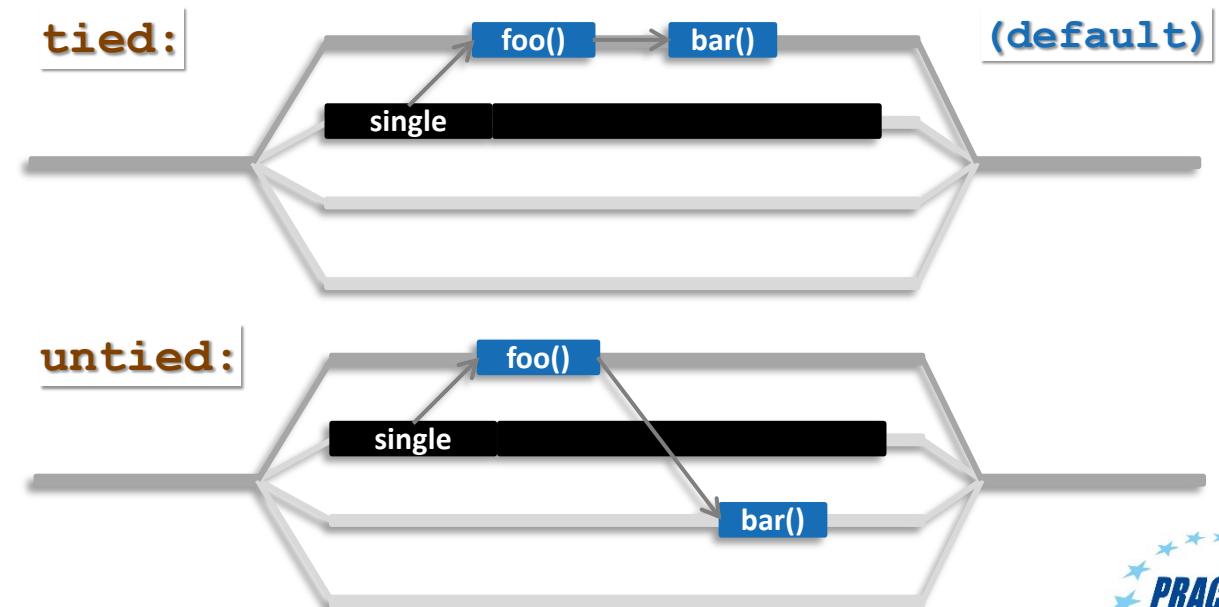
■ Task scheduling points (and the taskyield directive)

- tasks can be suspended/resumed at TSPs → some additional constraints to avoid deadlock problems
- implicit scheduling points (creation, synchronization, ...)
- explicit scheduling point: the taskyield directive

```
#pragma omp taskyield
```

■ Scheduling [tied/untied] tasks: example

```
#pragma omp parallel
#pragma omp single
{
    #pragma omp task untied
    {
        foo();
        #pragma omp taskyield
        bar()
    }
}
```



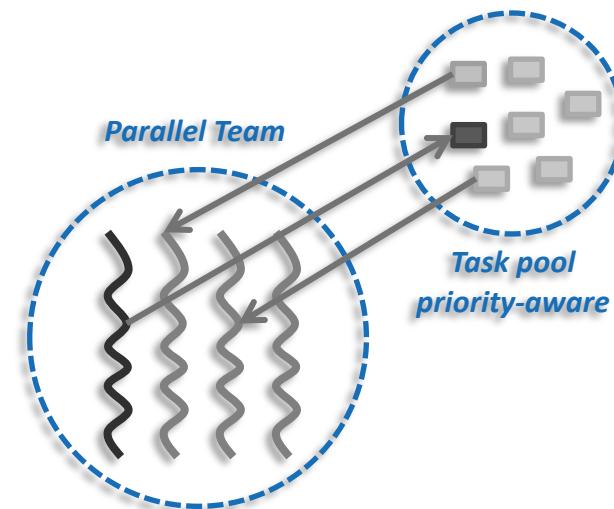
Task scheduling: programmer's hints

- Programmers may specify a priority value when creating a task

```
#pragma omp task priority(pvalue)
{structured-block}
```

- pvalue: the higher → the best (will be scheduled earlier)
- once a thread becomes idle, gets one of the highest priority tasks

```
#pragma omp parallel
#pragma omp single
{
    for ( i = 0; i < SIZE; i++) {
        #pragma omp task priority(1)
        { code_A; }
    }
    #pragma omp task priority(100)
    { code_B; }
    ...
}
```



Task synchronization: taskwait directive

- The taskwait directive (shallow task synchronization)

→ It is a stand-alone directive

```
#pragma omp taskwait
```

→ wait on the completion of child tasks of the current task; just direct children, not all descendant tasks;
includes an implicit task scheduling point (TSP)

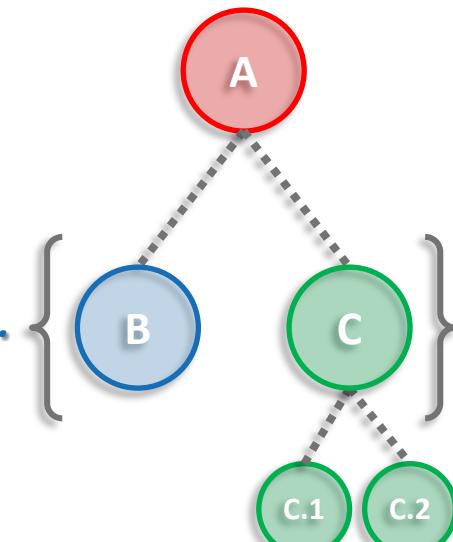
```
#pragma omp parallel
#pragma omp single
{
    #pragma omp task
    {
        #pragma omp task :B
        {
            #pragma omp task :C
            {
                ...
                #pragma omp task :C.x
            }
        }
        #pragma omp taskwait
    }
} // implicit barrier will wait for C.x
```

:A

:B

:C

wait for...



Task synchronization: barrier semantics

■ OpenMP barrier (implicit or explicit)

→ All tasks created by any thread of the current team are guaranteed to be completed at barrier exit

```
#pragma omp barrier
```

→ And all other implicit barriers at parallel, sections, for, single, etc...



Task synchronization: taskgroup construct

The taskgroup construct (deep task synchronization)

→ attached to a structured block; completion of all descendants of the current task; TSP at the end

```
#pragma omp taskgroup [clause[, clause]...]
{structured-block}
```

→ where clause (could only be): reduction(reduction-identifier: list-items)

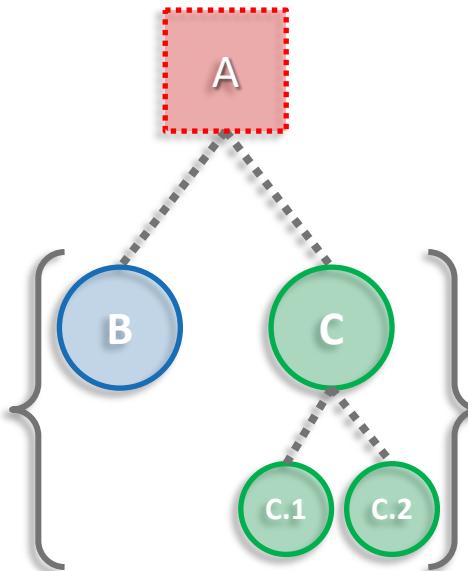
```
#pragma omp parallel
#pragma omp single
{
    #pragma omp taskgroup
    {
        #pragma omp task :A
        {
            #pragma omp task :B
            {
                ...
            }
            #pragma omp task :C
            {
                ...
                #C.1; #C.2; ...
            }
        } // end of taskgroup
    }
}
```

:A

:B

:C

wait for...



Data Environment



Explicit data-sharing clauses

■ Explicit data-sharing clauses (shared, private and firstprivate)

```
#pragma omp task shared(a)
{
    // Scope of a: shared
}
```

```
#pragma omp task private(b)
{
    // Scope of b: private
}
```

```
#pragma omp task firstprivate(c)
{
    // Scope of c: firstprivate
}
```

■ If **default** clause present, what the clause says

- shared: data which is not explicitly included in any other data sharing clause will be **shared**
- none: compiler will issue an error if the attribute is not explicitly set by the programmer (very useful!!!)

```
#pragma omp task default(shared)
{
    // Scope of all the references, not explicitly
    // included in any other data sharing clause,
    // and with no pre-determined attribute: shared
}
```

```
#pragma omp task default(none)
{
    // Compiler will force to specify the scope for
    // every single variable referenced in the context
}
```

Hint: Use default(none) to be forced to think about every variable if you do not see clearly.



Pre-determined data-sharing attributes

- threadprivate variables are threadprivate (1)
- dynamic storage duration objects are shared (malloc, new,...) (2)
- static data members are shared (3)
- variables declared inside the construct
 - static storage duration variables are shared (4)
 - automatic storage duration variables are private (5)
- the loop iteration variable(s)...

```
int A[SIZE];
#pragma omp threadprivate(A)
// ...
#pragma omp task
{
  // A: threadprivate
}
```

1

```
int *p;
p = malloc(sizeof(float)*SIZE);

#pragma omp task
{
  // *p: shared
}
```

2

```
#pragma omp task
{
  int x = MN;
  // Scope of x: private
}
```

5

```
#pragma omp task
{
  static int y;
  // Scope of y: shared
}
```

4

```
void foo(void) {
  static int s = MN;
}

#pragma omp task
{
  foo(); // s@foo(): shared
}
```

3

Implicit data-sharing attributes (in-practice)

- Implicit data-sharing rules for the task region
 - the **shared** attribute is lexically inherited
 - in any other case the variable is **firstprivate**

- Pre-determined rules (could not change)
- Explicit data-sharing clauses (+ default)
- Implicit data-sharing rules

- (in-practice) variable values within the task:

- value of a: 1
- value of b: x // undefined (undefined in parallel)
- value of c: 3
- value of d: 4
- value of e: 5

```
int a = 1;
void foo() {
    int b = 2, c = 3;
    #pragma omp parallel private(b)
    {
        int d = 4;
        #pragma omp task
        {
            int e = 5;
            // Scope of a:
            // Scope of b:
            // Scope of c:
            // Scope of d:
            // Scope of e:
        }
    }
}
```



Task reductions (using taskgroup)

■ Reduction operation

- perform some forms of recurrence calculations
- associative and commutative operators

■ The (taskgroup) scoping reduction clause

```
#pragma omp taskgroup task_reduction(op: list)
{structured-block}
```

- Register a new reduction at [1]
- Computes the final result after [3]

■ The (task) in_reduction clause [participating]

```
#pragma omp task in_reduction(op: list)
{structured-block}
```

- Task participates in a reduction operation [2]

```
int res = 0;
node_t* node = NULL;

...
#pragma omp parallel
{
    #pragma omp single
    {
        #pragma omp taskgroup task_reduction(+: res)
        { // [1]
            while (node) {
                #pragma omp task in_reduction(+: res) \
                    firstprivate(node)
                { // [2]
                    res += node->value;
                }
                node = node->next;
            }
        } // [3]
    }
}
```



Task reductions (+ modifiers)

■ Reduction modifiers

- Former reductions clauses have been extended
- task modifier allows to express task reductions
- Registering a new task reduction [1]
- Implicit tasks participate in the reduction [2]
- Compute final result after [4]

■ The (task) in_reduction clause [participating]

```
#pragma omp task in_reduction(op: list)
{structured-block}
```

- Task participates in a reduction operation [3]

```
int res = 0;
node_t* node = NULL;
...
#pragma omp parallel reduction(task,+: res)
{ // [1][2]
#pragma omp single
{
#pragma omp taskgroup
{
while (node) {
#pragma omp task in_reduction(+: res) \
firstprivate(node)
{ // [3]
res += node->value;
}
node = node->next;
}
}
} // [4]
```



Tasking illustrated



Fibonacci illustrated

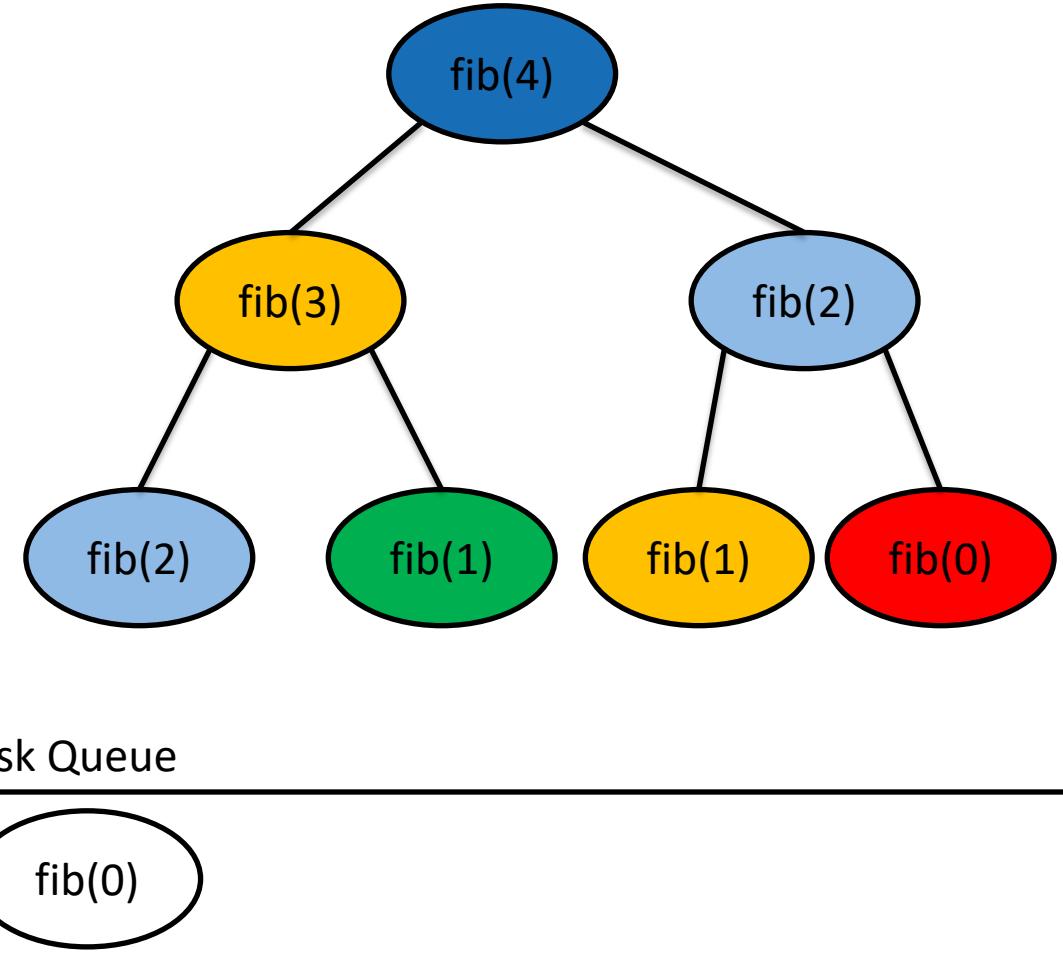
```
1 int main(int argc,
2          char* argv[])
3 {
4     [...]
5     #pragma omp parallel
6     {
7         #pragma omp single
8         {
9             fib(input);
10        }
11    }
12    [...]
13 }
```

```
14 int fib(int n)  {
15     if (n < 2) return n;
16     int x, y;
17     #pragma omp task shared(x)
18     {
19         x = fib(n - 1);
20     }
21     #pragma omp task shared(y)
22     {
23         y = fib(n - 2);
24     }
25     #pragma omp taskwait
26     return x+y;
27 }
```

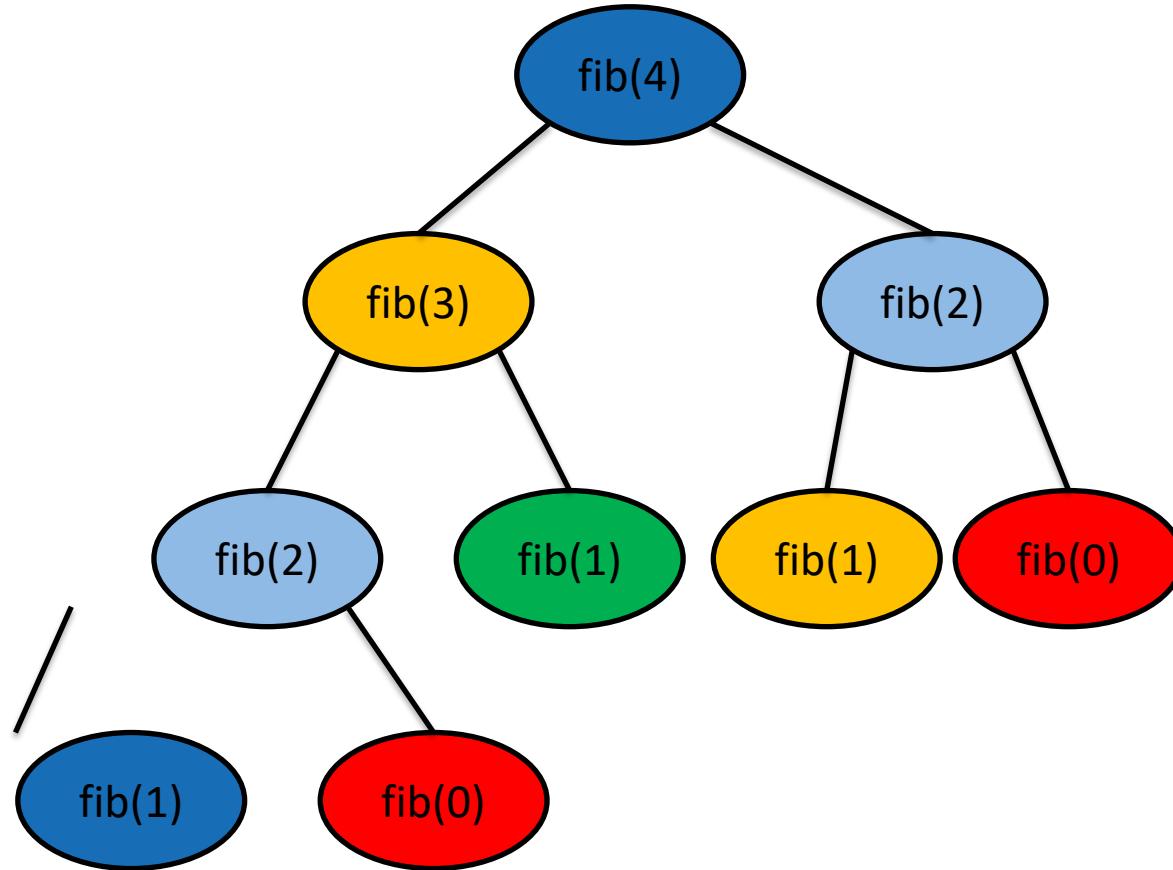
- Only one Task / Thread enters fib() from main(), it is responsible for creating the two initial work tasks
- Taskwait is required, as otherwise x and y would get lost



- T1 enters fib(4)
- T1 creates tasks for fib(3) and fib(2)
- T1 and T2 execute tasks from the queue
- T1 and T2 create 4 new tasks
- T1 - T4 execute tasks



- T1 enters fib(4)
- T1 creates tasks for fib(3) and fib(2)
- T1 and T2 execute tasks from the queue
- T1 and T2 create 4 new tasks
- T1 - T4 execute tasks
- ...



The taskloop Construct



Traditional Worksharing

- Worksharing constructs do not compose well
- Pathological example: parallel dgemm in MKL

```
void example() {
    #pragma omp parallel
    {
        compute_in_parallel(A);
        compute_in_parallel_too(B);
        // dgemm is either parallel or sequential,
        // but has no orphaned worksharing
        cblas_dgemm(CblasRowMajor, CblasNoTrans, CblasNoTrans,
                    m, n, k, alpha, A, k, B, n, beta, C, n);

    }
}
```

- Writing such code either
 - oversubscribes the system,
 - yields bad performance due to OpenMP overheads, or
 - needs a lot of glue code to use sequential dgemm only for sub-matrixes



Example: Sparse CG

```

for (iter = 0; iter < sc->maxIter; iter++) {
    precon(A, r, z);
    vectorDot(r, z, n, &rho);
    beta = rho / rho_old;
    xpay(z, beta, n, p);
    matvec(A, p, q);
    vectorDot(p, q, n, &dot_pq);
    alpha = rho / dot_pq;
    axpy(alpha, p, n, x);
    axpy(-alpha, q, n, r);
    sc->residual = sqrt(rho) * b;
    if (sc->residual <= sc->toler)
        break;
    rho_old = rho;
}

```

```

void matvec(Matrix *A, double *x, double *y) {
    // ...
#pragma omp parallel for \
    private(i,j,is,ie,j0,y0) \
    schedule(static)
for (i = 0; i < A->n; i++) {
    y0 = 0;
    is = A->ptr[i];
    ie = A->ptr[i + 1];
    for (j = is; j < ie; j++) {
        j0 = index[j];
        y0 += value[j] * x[j0];
    }
    y[i] = y0;
}
// ...
}

```



Tasking use case: saxpy (taskloop)

```
for ( i = 0; i<SIZE; i+=1) {
    A[i]=A[i]*B[i]*S;
}
```

```
for ( i = 0; i<SIZE; i+=TS) {
    UB = SIZE < (i+TS)?SIZE:i+TS;
    for ( ii=i; ii<UB; ii++) {
        A[ii]=A[ii]*B[ii]*S;
    }
}
```

```
#pragma omp parallel
#pragma omp single
for ( i = 0; i<SIZE; i+=TS) {
    UB = SIZE < (i+TS)?SIZE:i+TS;
    #pragma omp task private(ii) \
        firstprivate(i,UB) shared(S,A,B)
    for ( ii=i; ii<UB; ii++) {
        A[ii]=A[ii]*B[ii]*S;
    }
}
```

- Difficult to determine grain
 - 1 single iteration → too fine
 - whole loop → no parallelism
- Manually transform the code
 - blocking techniques
- Improving programmability
 - OpenMP taskloop

```
#pragma omp taskloop grainsize(TS)
for ( i = 0; i<SIZE; i+=1) {
    A[i]=A[i]*B[i]*S;
}
```

- Hiding the internal details
- Grain size ~ Tile size (TS) → but implementation decides exact grain size



The taskloop Construct

- Task generating construct: decompose a loop into chunks, create a task for each loop chunk

```
#pragma omp taskloop [clause[, clause]...]
{structured-for-loops}
```

```
!$omp taskloop [clause[, clause]...]
...structured-do-loops...
 !$omp end taskloop
```

- Where clause is one of:

- shared(list)
- private(list)
- firstprivate(list)
- lastprivate(list)
- default(sh | pr | fp | none)
- reduction(r-id: list)
- in_reduction(r-id: list)

Data Environment

- grainsize(grain-size)
- num_tasks(num-tasks)

Chunks/Grain

- if(scalar-expression)
- final(scalar-expression)
- mergeable

Cutoff Strategies

- untied
- priority(priority-value)

Scheduler (R/H)

- collapse(n)
- nogroup
- allocate([allocator:] list)

Miscellaneous

Worksharing vs. taskloop constructs (1/2)

```
subroutine worksharing
    integer :: x
    integer :: i
    integer, parameter :: T = 16
    integer, parameter :: N = 1024

    x = 0
!$omp parallel shared(x) num_threads(T)

!$omp do
    do i = 1,N
!$omp atomic
        x = x + 1
!$omp end atomic
    end do
!$omp end do

!$omp end parallel
    write (*,'(A,I0)') 'x = ', x
end subroutine
```

Result: x = 1024

```
subroutine taskloop
    integer :: x
    integer :: i
    integer, parameter :: T = 16
    integer, parameter :: N = 1024

    x = 0
!$omp parallel shared(x) num_threads(T)

!$omp taskloop
    do i = 1,N
!$omp atomic
        x = x + 1
!$omp end atomic
    end do
!$omp end taskloop

!$omp end parallel
    write (*,'(A,I0)') 'x = ', x
end subroutine
```

Result: x = 16384

Worksharing vs. taskloop constructs (2/2)

```
subroutine worksharing
    integer :: x
    integer :: i
    integer, parameter :: T = 16
    integer, parameter :: N = 1024

    x = 0
!$omp parallel shared(x) num_threads(T)

!$omp do
    do i = 1,N
!$omp atomic
        x = x + 1
!$omp end atomic
    end do
!$omp end do

!$omp end parallel
    write (*,'(A,I0)') 'x = ', x
end subroutine
```

Result: x = 1024

```
subroutine taskloop
    integer :: x
    integer :: i
    integer, parameter :: T = 16
    integer, parameter :: N = 1024

    x = 0
!$omp parallel shared(x) num_threads(T)
!$omp single
!$omp taskloop
    do i = 1,N
!$omp atomic
        x = x + 1
!$omp end atomic
    end do
!$omp end taskloop
!$omp end single
!$omp end parallel
    write (*,'(A,I0)') 'x = ', x
end subroutine
```

Result: x = 1024

Taskloop decomposition approaches

■ Clause: grainsize(grain-size)

- Chunks have at least grain-size iterations
- Chunks have maximum 2x grain-size iterations

```
int TS = 4 * 1024;  
#pragma omp taskloop grainsize(TS)  
for ( i = 0; i<SIZE; i+=1) {  
    A[i]=A[i]*B[i]*S;  
}
```

■ Clause: num_tasks(num-tasks)

- Create num-tasks chunks
- Each chunk must have at least one iteration

```
int NT = 4 * omp_get_num_threads();  
#pragma omp taskloop num_tasks(NT)  
for ( i = 0; i<SIZE; i+=1) {  
    A[i]=A[i]*B[i]*S;  
}
```

■ If none of previous clauses is present, the *number of chunks* and the *number of iterations per chunk* is implementation defined

■ Additional considerations:

- The order of the creation of the loop tasks is unspecified
- Taskloop creates an implicit taskgroup region; **nogroup** → no implicit taskgroup region is created



Collapsing iteration spaces with taskloop

■ The collapse clause in the taskloop construct

```
#pragma omp taskloop collapse(n)
{structured-for-loops}
```

- Number of loops associated with the taskloop construct (n)
- Loops are collapsed into one larger iteration space
- Then divided according to the **grainsize** and **num_tasks**

■ Intervening code between any two associated loops

- at least once per iteration of the enclosing loop
- at most once per iteration of the innermost loop

```
#pragma omp taskloop collapse(2)
for ( i = 0; i<SX; i+=1) {
    for ( j= 0; i<SY; j+=1) {
        for ( k = 0; i<SZ; k+=1) {
            A[f(i,j,k)]=<expression>;
        }
    }
}
```



```
#pragma omp taskloop
for ( ij = 0; i<SX*SY; ij+=1) {
    for ( k = 0; i<SZ; k+=1) {
        i = index_for_i(ij);
        j = index_for_j(ij);
        A[f(i,j,k)]=<expression>;
    }
}
```

Task reductions (using taskloop)

Clause: reduction(r-id: list)

- It defines the scope of a new reduction
- All created tasks participate in the reduction
- It cannot be used with the **nogroup** clause

```
double dotprod(int n, double *x, double *y) {
    double r = 0.0;
#pragma omp taskloop reduction(+: r)
    for (i = 0; i < n; i++)
        r += x[i] * y[i];

    return r;
}
```

Clause: in_reduction(r-id: list)

- Reuse an already defined reduction scope
- All created tasks participate in the reduction
- It can be used with the **nogroup*** clause, but it is user responsibility to guarantee result

```
double dotprod(int n, double *x, double *y) {
    double r = 0.0;
#pragma omp taskgroup task_reduction(+: r)
{
    #pragma omp taskloop in_reduction(+: r)*
    for (i = 0; i < n; i++)
        r += x[i] * y[i];
}
return r;
}
```



Composite construct: taskloop simd

- Task generating construct: decompose a loop into chunks, create a task for each loop chunk
- Each generated task will apply (internally) SIMD to each loop chunk
 - C/C++ syntax:

```
#pragma omp taskloop simd [clause[,] clause]...
{structured-for-loops}
```

- Fortran syntax:

```
!$omp taskloop simd [clause[,] clause]...
...structured-do-loops...
 !$omp end taskloop
```

- Where clause is any of the clauses accepted by **taskloop** or **simd** directives



Improving Tasking Performance: Task dependences



■ Task dependences as a way to define task-execution constraints

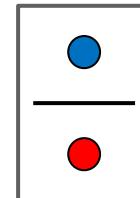
```
int x = 0;  
#pragma omp parallel  
#pragma omp single  
{  
    #pragma omp task  
    std::cout << x << std::endl;  
  
    #pragma omp taskwait  
  
    #pragma omp task  
    x++;  
}
```

OpenMP 3.1

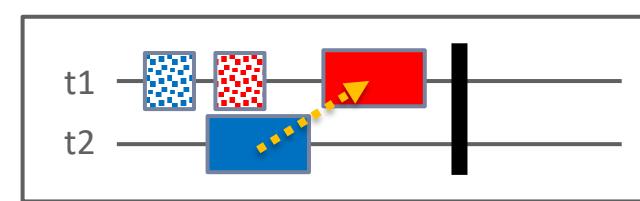
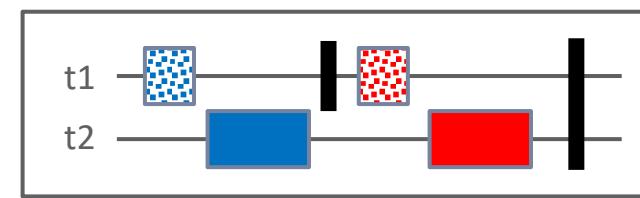
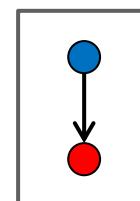
```
int x = 0;  
#pragma omp parallel  
#pragma omp single  
{  
    #pragma omp task depend(in: x)  
    std::cout << x << std::endl;  
  
    #pragma omp task depend(inout: x)  
    x++;  
}
```

OpenMP 4.0

OpenMP 3.1



OpenMP 4.0



Task's creation time
 Task's execution time



■ Task dependences as a way to define task-execution constraints

```
int x = 0;  
#pragma omp parallel  
#pragma omp single  
{  
    #pragma omp task  
    std::cout << x << std::endl;  
  
    #pragma omp taskwait  
  
    #pragma omp task  
    x++;  
}
```

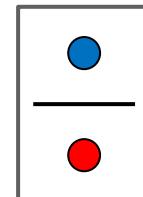
OpenMP 3.1

```
int x = 0;  
#pragma omp parallel  
#pragma omp single  
{  
    #pragma omp task depend(in: x)  
    std::cout << x << std::endl;  
  
    x++;  
}  
  
end(inout: x)
```

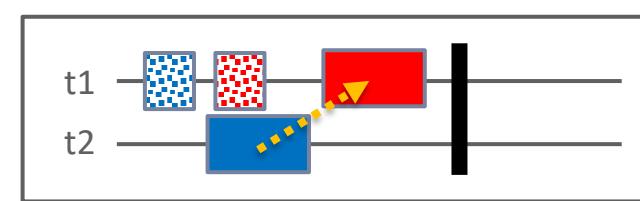
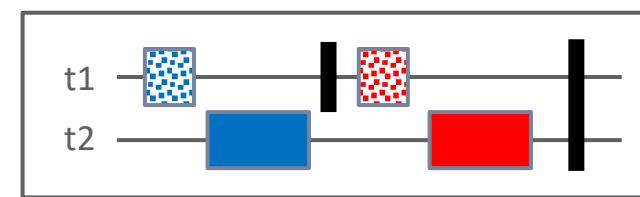
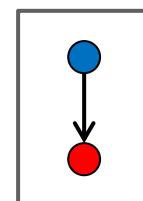
OpenMP 4.0

Task dependences can help us to remove
“strong” synchronizations, increasing the look
ahead and, frequently, the parallelism!!!!

OpenMP 3.1



OpenMP 4.0



Task's creation time
Task's execution time

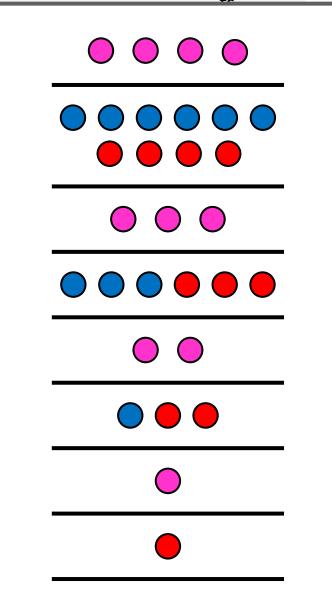


Motivation: Cholesky factorization

```
void cholesky(int ts, int nt, double* a[nt][nt]) {
    for (int k = 0; k < nt; k++) {
        // Diagonal Block factorization
        potrf(a[k][k], ts, ts);

        // Triangular systems
        for (int i = k + 1; i < nt; i++) {
            #pragma omp task
            trsm(a[k][k], a[k][i], ts, ts);
        }
        #pragma omp taskwait

        // Update trailing matrix
        for (int i = k + 1; i < nt; i++) {
            for (int j = k + 1; j < i; j++)
                #pragma omp task
                dgemm(a[k][i], a[k][j], a[j][i], ts, ts,
                );
            #pragma omp task
            syrk(a[k][i], a[i][i], ts, ts);
        }
        #pragma omp taskwait
    }
}
```

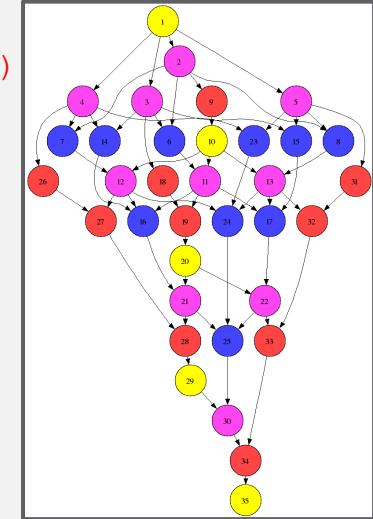


OpenMP 3.1

```
void cholesky(int ts, int nt, double* a[nt][nt]) {
    for (int k = 0; k < nt; k++) {
        // Diagonal Block factorization
        #pragma omp task depend(inout: a[k][k])
        potrf(a[k][k], ts, ts);

        // Triangular systems
        for (int i = k + 1; i < nt; i++) {
            #pragma omp task depend(in: a[k][k])
            #pragma omp task depend(inout: a[k][i])
            trsm(a[k][k], a[k][i], ts, ts);
        }

        // Update trailing matrix
        for (int i = k + 1; i < nt; i++) {
            for (int j = k + 1; j < i; j++)
                #pragma omp task depend(inout: a[j][i])
                #pragma omp task depend(in: a[k][i], a[k][j])
                dgemm(a[k][i], a[k][j], a[j][i], ts, ts,
                );
            #pragma omp task depend(inout: a[i][i])
            #pragma omp task depend(in: a[k][i])
            syrk(a[k][i], a[i][i], ts, ts);
        }
    }
}
```



OpenMP 4.0

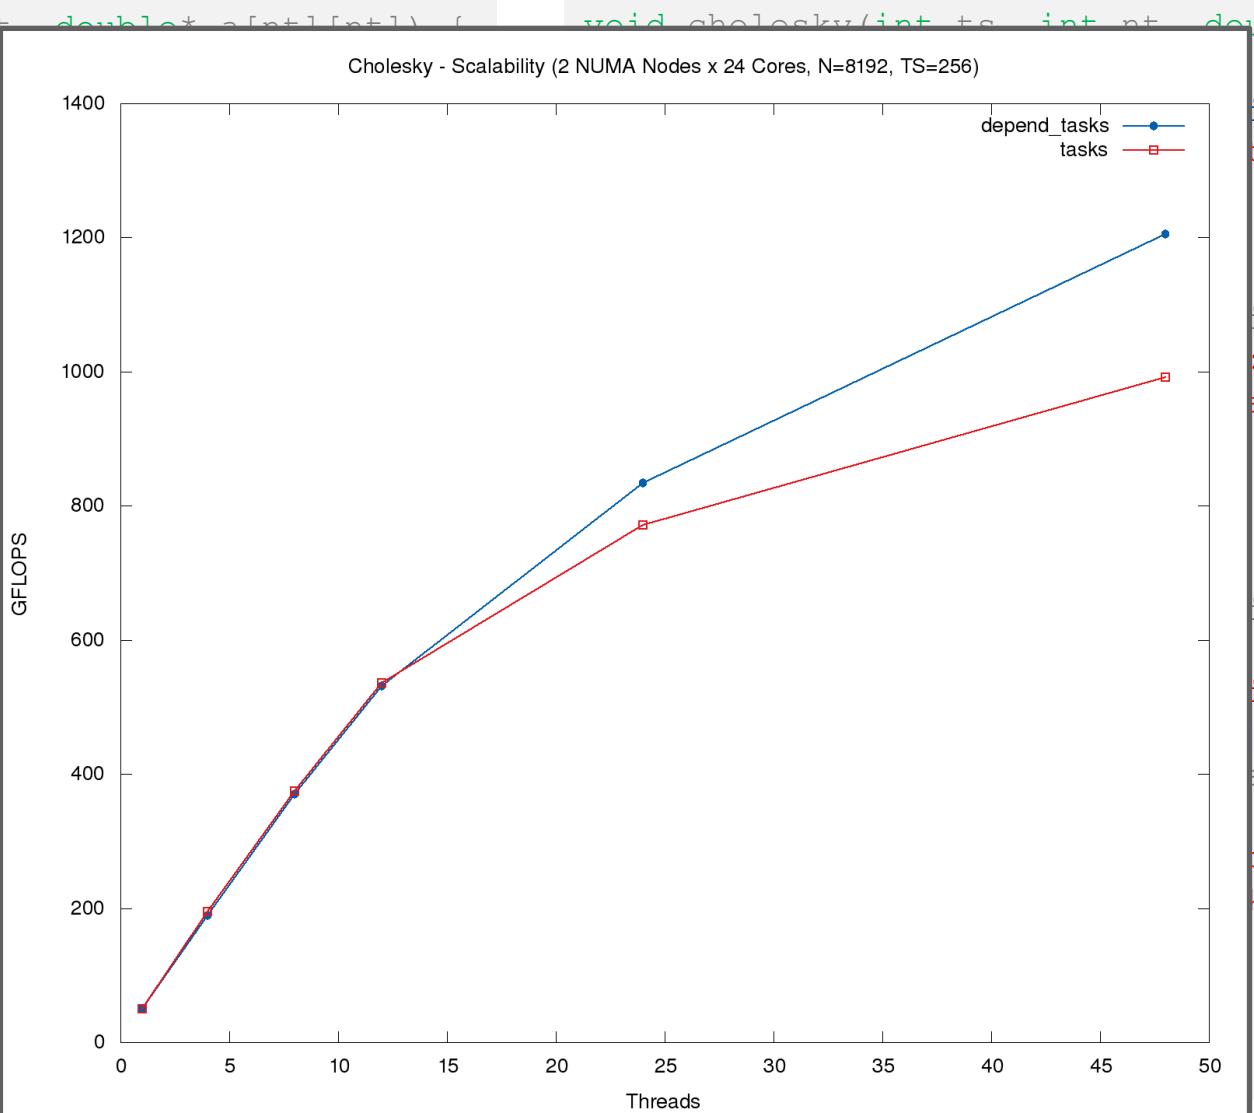


Motivation: Cholesky factorization

```
void cholesky(int ts, int nt, double* a[nt][nt]) {
    for (int k = 0; k < nt; k++) {
        // Diagonal Block factorization
        potrf(a[k][k], ts, ts);

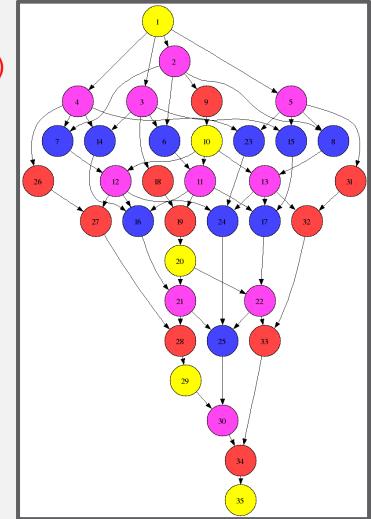
        // Triangular systems
        for (int i = k + 1; i < nt; i++) {
            #pragma omp task
            trsm(a[k][k], a[k][i]);
        }
        #pragma omp taskwait

        // Update trailing matrix
        for (int i = k + 1; i < nt; i++) {
            for (int j = k + 1; j < i; j++) {
                #pragma omp task
                dgemm(a[k][i], a[k][j], 1.0);
            }
            #pragma omp task
            syrk(a[k][i], a[i][i], ts, ts);
        }
        #pragma omp taskwait
    }
}
```



```
    ion
    t: a[k][k])
    i++) {
    j++) {
    inout: a[j][i])
    [k][i], a[k][j])
    a[j][i], ts, ts);

    out: a[i][i])
    <][i])
    ts);
```



OpenMP 4.0



What's in the spec



What's in the spec: a bit of history

OpenMP 4.0

- The **depend clause** was added to the **task construct**

OpenMP 4.5

- The **depend clause** was added to the **target constructs**
- Support to **doacross loops**

OpenMP 5.0

- **lvalue expressions** in the **depend clause**
- New dependency type: **mutexinoutset**
- Iterators were added to the **depend clause**
- The **depend clause** was added to the **taskwait construct**
- **Dependable objects**



What's in the spec: syntax depend clause

```
depend( [depend-modifier,] dependency-type: list-items)
```

where:

- depend-modifier **is used to define iterators**
- dependency-type **may be:** in, out, inout, mutexinoutset and depobj
- A list-item **may be:**
 - C/C++: A lvalue expr or an array section `depend(in: x, v[i], *p, w[10:10])`
 - Fortran: A variable or an array section `depend(in: x, v(i), w(10:20))`



What's in the spec: sema depend clause (1)

- A task cannot be executed until all its predecessor tasks are completed
- If a task defines an `in` dependence over a list-item
 - the task will depend on all previously generated sibling tasks that reference that list-item in an `out` or `inout` dependence
- If a task defines an `out/inout` dependence over list-item
 - the task will depend on all previously generated sibling tasks that reference that list-item in an `in`, `out` or `inout` dependence

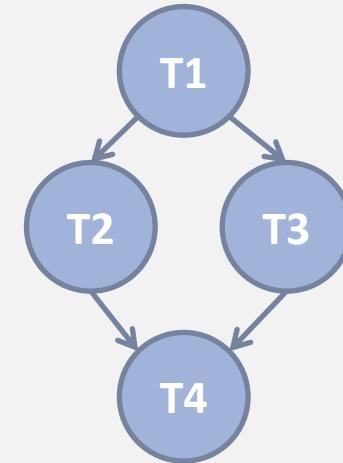
What's in the spec: depend clause (1)

- A task cannot be executed until all its predecessor tasks are completed

- If a task defines

→ the task will complete
an out or in

```
int x = 0;  
#pragma omp parallel  
#pragma omp single  
{  
    #pragma omp task depend(inout: x) //T1  
    { ... }  
  
    #pragma omp task depend(in: x)      //T2  
    { ... }  
  
    #pragma omp task depend(in: x)      //T3  
    { ... }  
  
    #pragma omp task depend(inout: x) //T4  
    { ... }  
}
```



- If a task defines

→ the task will complete
an in, out or inout

none of the list items in

none of the list items in

What's in the spec: depend clause (2)

■ New dependency type: mutexinoutset

```

int x = 0, y = 0, res = 0;
#pragma omp parallel
#pragma omp single
{
    #pragma omp task depend(out: res)    //T0
    res = 0;

    #pragma omp task depend(out: x)      //T1
    long_computation(x);

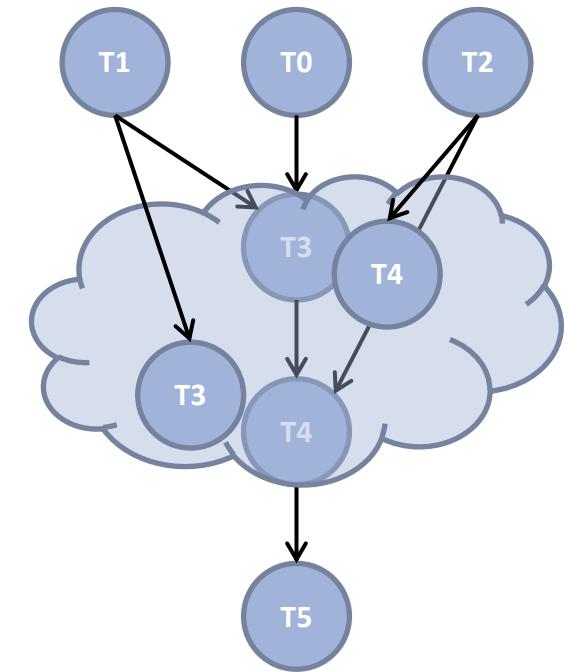
    #pragma omp task depend(out: y)      //T2
    short_computation(y);

    #pragma omp task depend(in: x) depend(mutexinoutset:T3res) //T3
    res += x;

    #pragma omp task depend(in: y) depend(mutexinoutset:T4res) //T4
    res += y;

    #pragma omp task depend(in: res)    //T5
    std::cout << res << std::endl;
}

```



1. *inoutset property*: tasks with a `mutexinoutset` dependence create a cloud of tasks (an inout set) that synchronizes with previous & posterior tasks that dependent on the same list item
2. *mutex property*: Tasks inside the inout set can be executed in any order but with mutual exclusion

What's in the spec: depend clause (4)

- Task dependences are defined among **sibling tasks**

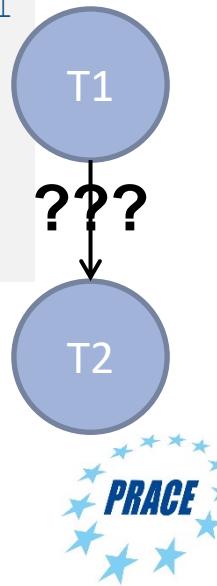
- List items used in the depend clauses [...] must indicate **identical** or **disjoint storage**

```
//test1.cc
int x = 0;
#pragma omp parallel
#pragma omp single
{
    #pragma omp task depend(inout: x)    //T1
    {
        #pragma omp task depend(inout: x) //T1.1
        x++;

        #pragma omp taskwait
    }
    #pragma omp task depend(in: x) //T2
    std::cout << x << std::endl;
}
```

```
//test2.cc
int a[100] = {0};
#pragma omp parallel
#pragma omp single
{
    #pragma omp task depend(inout: a[50:99]) //T1
    compute(/* from */ &a[50], /*elems*/ 50);

    #pragma omp task depend(in: a)    //T2
    print(/* from */ a, /* elem */ 100);
}
```



What's in the spec: depend clause (5)

■ Iterators + deps: a way to define a dynamic number of dependences

```
std::list<int> list = ...;
int n = list.size();

#pragma omp parallel
#pragma omp single
{
    for (int i = 0; i < n; ++i)
        #pragma omp task depend(out: list[i])           //Px
        compute_elem(list[i]);

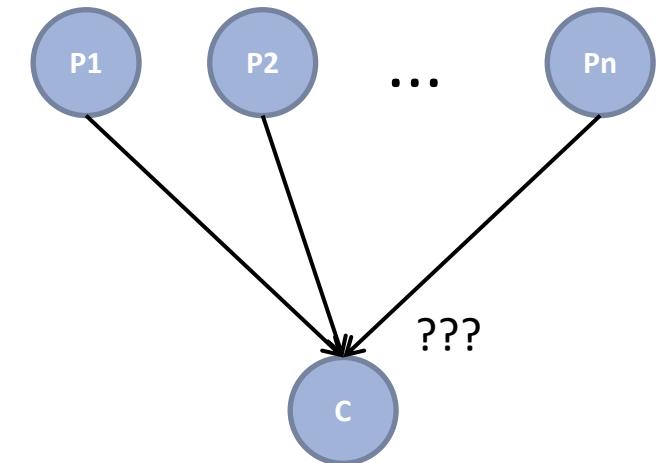
    #pragma omp task depend(iterator(j=0:n), in : list[j]) //C
    print_elems(list);
}
```

It seems innocent but it's not:
`depend(out: list.operator[](i))`

`//Px`

`//C`

Equivalent to:
`depend(in: list[0], list[1], ..., list[n-1])`



Use case

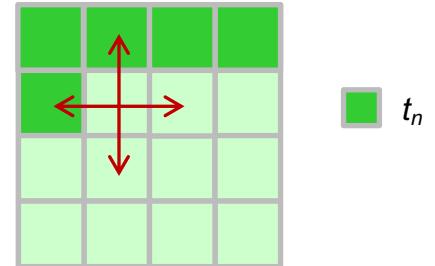


Use case: intro to Gauss-seidel

```
void serial_gauss_seidel(int tsteps, int size, int (*p) [size]) {
    for (int t = 0; t < tsteps; ++t) {
        for (int i = 1; i < size-1; ++i) {
            for (int j = 1; j < size-1; ++j) {
                p[i][j] = 0.25 * (p[i][j-1] + // left
                                    p[i][j+1] + // right
                                    p[i-1][j] + // top
                                    p[i+1][j]); // bottom
            }
        }
    }
}
```

Access pattern analysis

For a specific t, i and j



Each cell depends on:

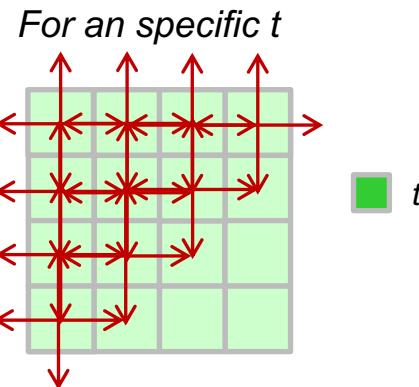
- two cells (north & west) that are computed in the current time step, and
- two cells (south & east) that were computed in the previous time step



Use case: Gauss-seidel (2)

```
void serial_gauss_seidel(int tsteps, int size, int (*p) [size]) {
    for (int t = 0; t < tsteps; ++t) {
        for (int i = 1; i < size-1; ++i) {
            for (int j = 1; j < size-1; ++j) {
                p[i][j] = 0.25 * (p[i][j-1] + // left
                                    p[i][j+1] + // right
                                    p[i-1][j] + // top
                                    p[i+1][j]); // bottom
            }
        }
    }
}
```

1st parallelization strategy



We can exploit the **wavefront** to obtain parallelism!!



Use case : Gauss-seidel (3)

```
void gauss_seidel(int tsteps, int size, int TS, int (*p)[size]) {
    int NB = size / TS;
    #pragma omp parallel
    for (int t = 0; t < tsteps; ++t) {
        // First NB diagonals
        for (int diag = 0; diag < NB; ++diag) {
            #pragma omp for
            for (int d = 0; d <= diag; ++d) {
                int ii = d;
                int jj = diag - d;
                for (int i = 1+ii*TS; i < ((ii+1)*TS); ++i)
                    for (int j = 1+jj*TS; i < ((jj+1)*TS); ++j)
                        p[i][j] = 0.25 * (p[i][j-1] + p[i][j+1] +
                                           p[i-1][j] + p[i+1][j]);
            }
        }
        // Lasts NB diagonals
        for (int diag = NB-1; diag >= 0; --diag) {
            // Similar code to the previous loop
        }
    }
}
```

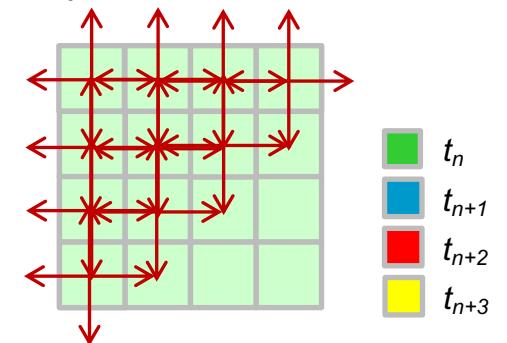


Use case : Gauss-seidel (4)

```
void serial_gauss_seidel(int tsteps, int size, int (*p) [size]) {
    for (int t = 0; t < tsteps; ++t) {
        for (int i = 1; i < size-1; ++i) {
            for (int j = 1; j < size-1; ++j) {
                p[i][j] = 0.25 * (p[i][j-1] + // left
                                    p[i][j+1] + // right
                                    p[i-1][j] + // top
                                    p[i+1][j]); // bottom
            }
        }
    }
}
```

2nd parallelization strategy

multiple time iterations



We can exploit the **wavefront** of multiple time steps to obtain **MORE** parallelism!!



Use case : Gauss-seidel (5)

```

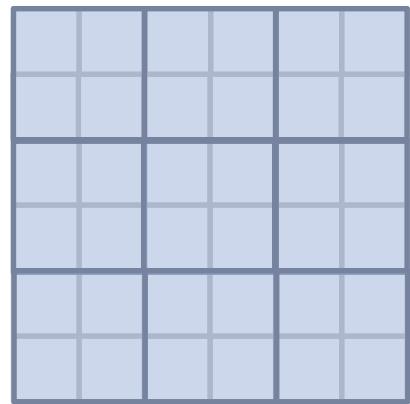
void gauss_seidel(int tsteps, int size, int TS, int (*p)[size]) {
    int NB = size / TS;

    #pragma omp parallel
    #pragma omp single
    for (int t = 0; t < tsteps; ++t)
        for (int ii=1; ii < size-1; ii+=TS)
            for (int jj=1; jj < size-1; jj+=TS) {
                #pragma omp task depend(inout: p[ii:TS][jj:TS])
                depend(in: p[ii-TS:TS][jj:TS], p[ii+TS:TS][jj:TS],
                       p[ii:TS][jj-TS:TS], p[ii:TS][jj+TS:TS])

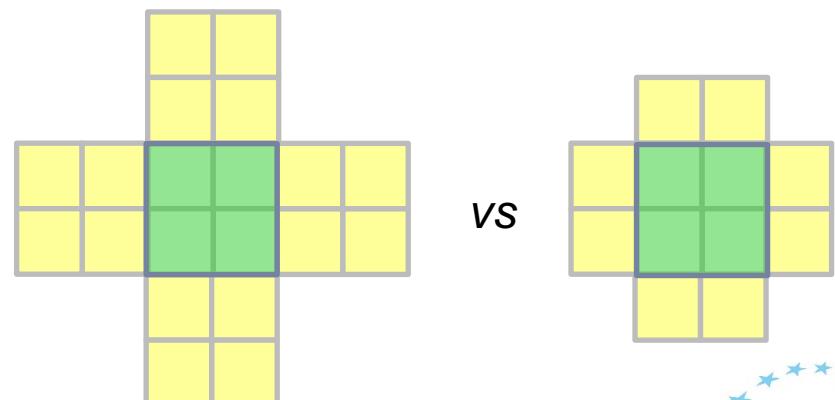
                {
                    for (int i=ii; i<(1+ii)*TS; ++i)
                        for (int j=jj; j<(1+jj)*TS; ++j)
                            p[i][j] = 0.25 * (p[i][j-1] + p[i][j+1] +
                                              p[i-1][j] + p[i+1][j]);
                }
            }
}

```

inner matrix region



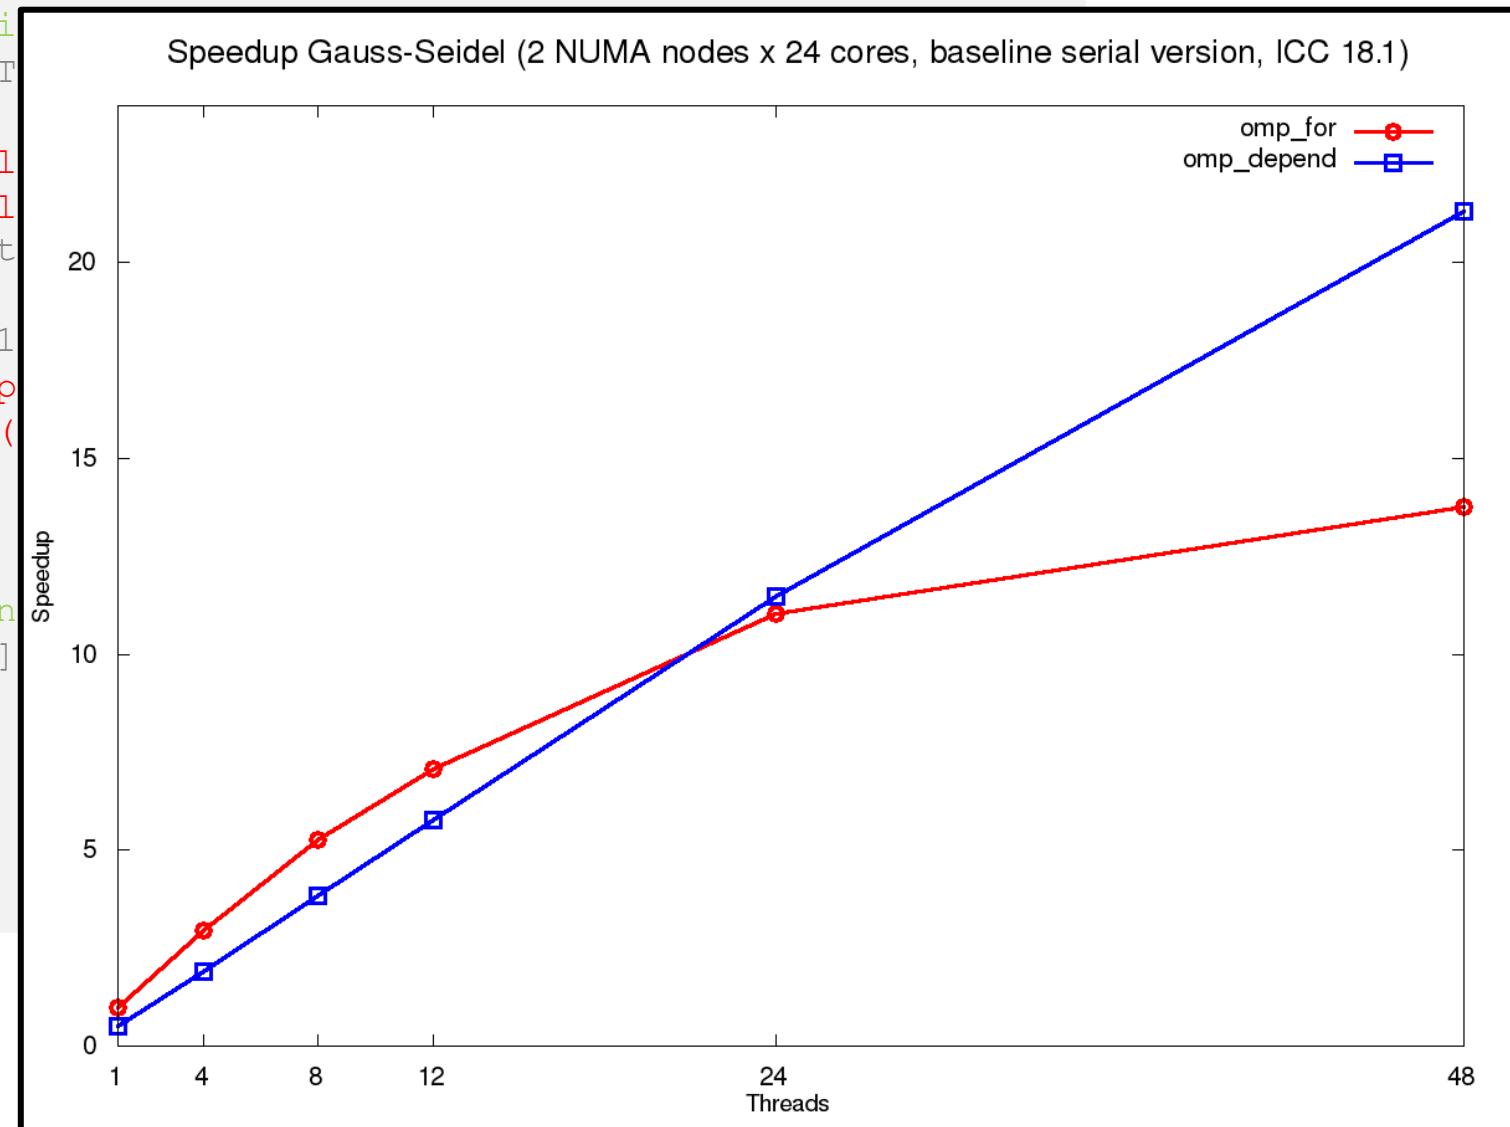
Q: Why do the input dependences depend on the whole block rather than just a column/row?



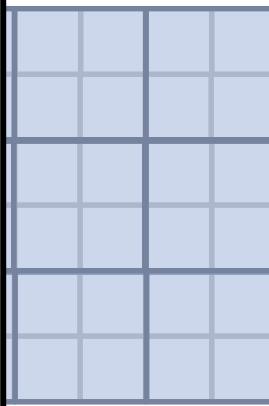
Use case : Gauss-seidel (5)

```
void gauss_seidel(i
    int NB = size / T

#pragma omp paral
#pragma omp singl
for (int t = 0; t
    for (int ii=1;
        for (int jj=1
            #pragma omp
            depend(
{
    for (int
        for (in
            p[i]
}
}
}
```

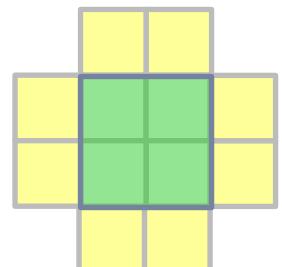


matrix region



the input dependences
use whole block rather
than a column/row?

vs



Improving Tasking Performance: Cutoff clauses and strategies



Example: Sudoku revisited



Parallel Brute-force Sudoku

- This parallel algorithm finds all valid solutions

	6					8	11			15	14				16
15	11			16	14			12			6				
13		9	12			3	16	14		15	11	10			
2	16		11	15	10	1									
	15	11	10		16	2	13	8	9	12					
12	13			4	1	5	6	2	3				11	10	
5		6	1	12		9	15	11	10	7	16				3
	2			10		11	6		5		13				9
10	7	15	11	16			12	13							6
9					1		2		16	10					11
1	4	6	9	13		7	11		3	16					
16	14		7	10	15	4	6	1			13	8			
11	10	15			16	9	12	13		1	5	4			
	12		1	4	6	16			11	10					
	5		8	12	13	10		11	2						14
3	16		10		7		6				12				

- (1) Search an empty field
- (2) Try all numbers:
 - (2 a) Check Sudoku
 - If invalid: skip
 - If valid: Go to next field
- Wait for completion

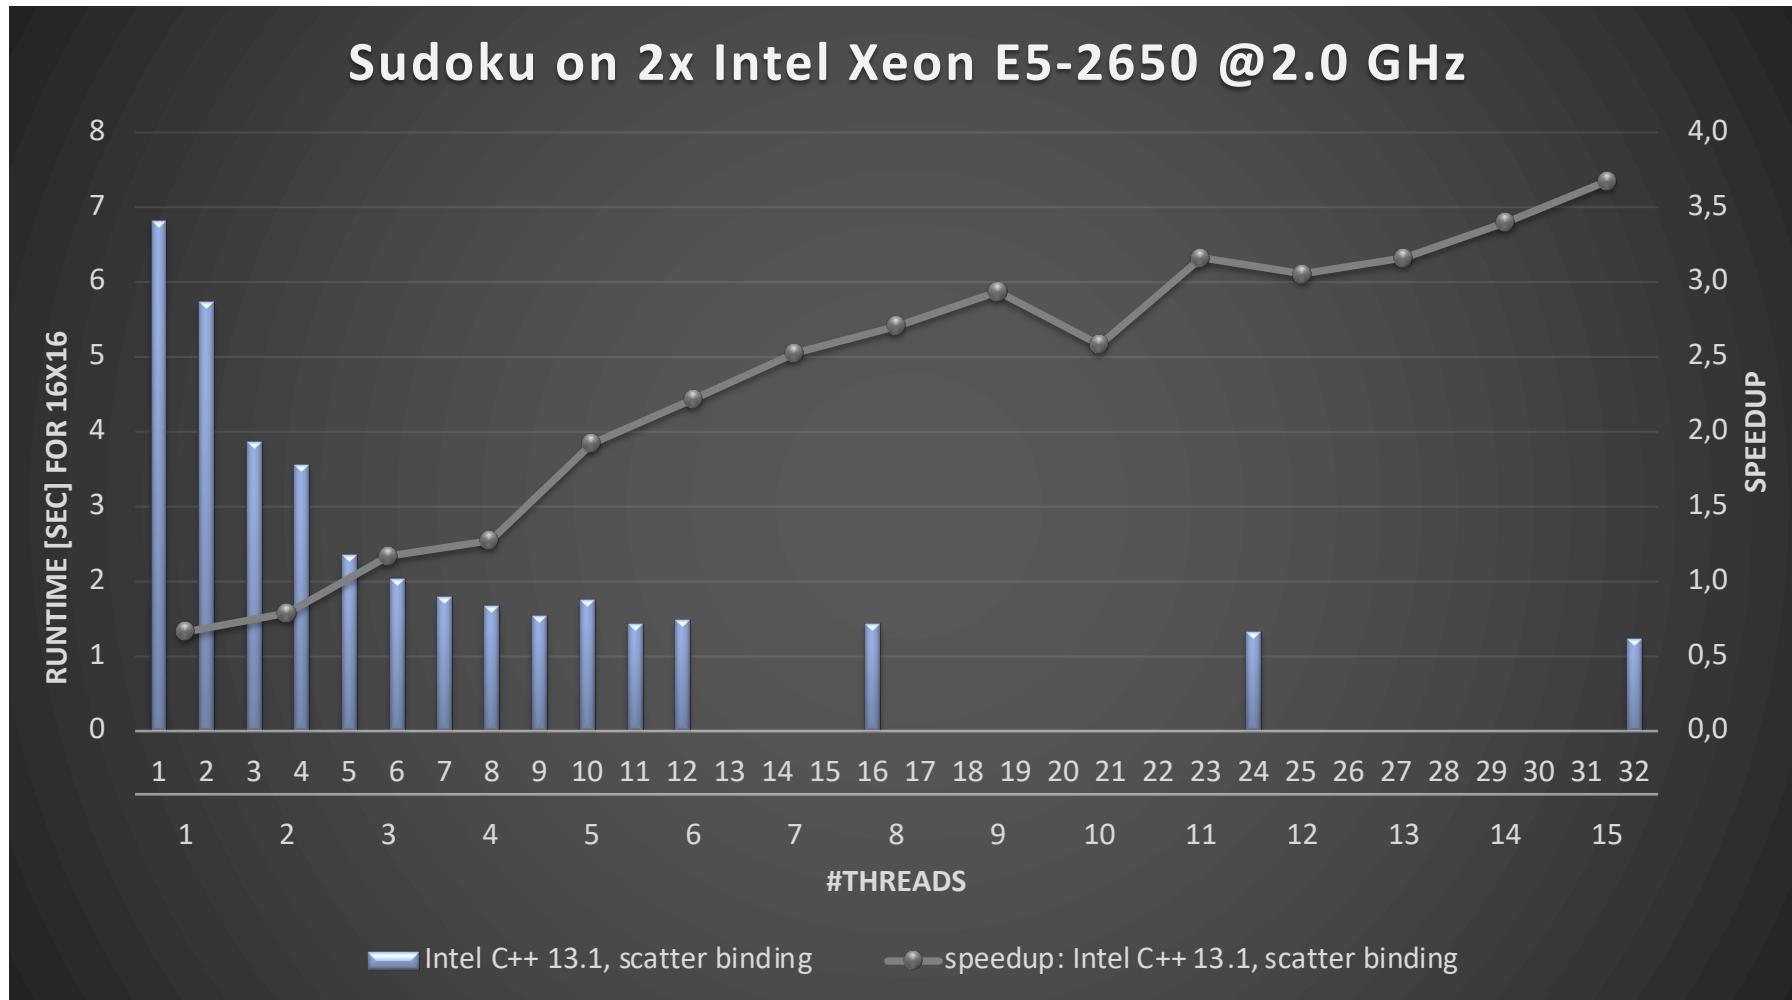
first call contained in a
`#pragma omp parallel`
`#pragma omp single`
 such that one tasks starts the
 execution of the algorithm

`#pragma omp task`
 needs to work on a new copy
 of the Sudoku board

`#pragma omp taskwait`
 wait for all child tasks

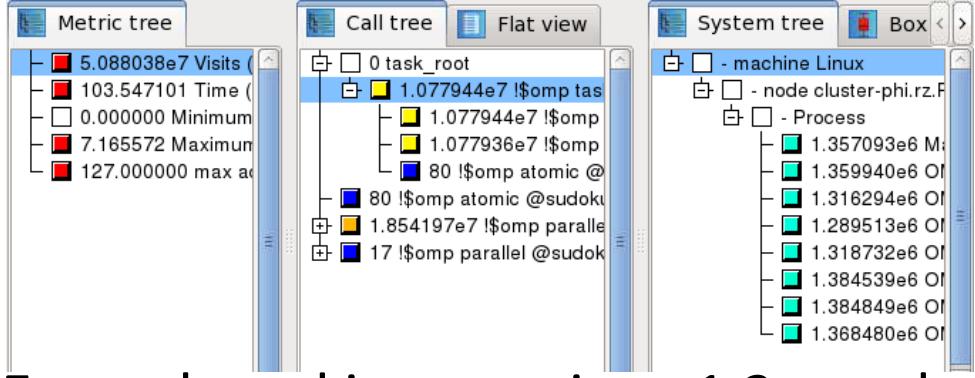


Performance Evaluation

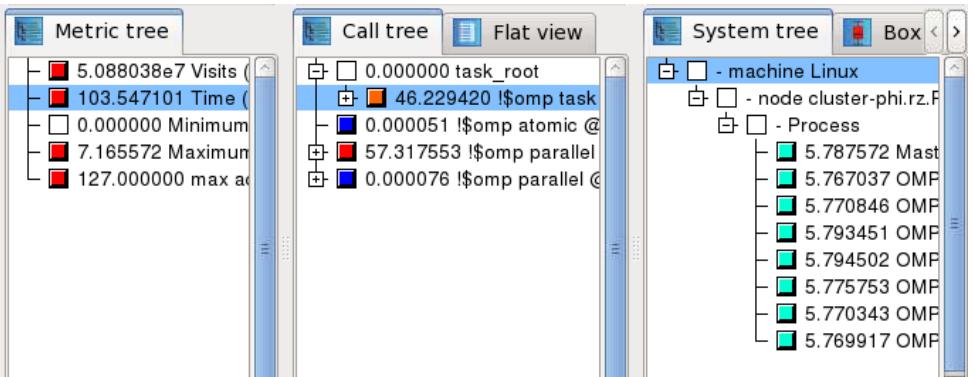


Performance Analysis

Event-based profiling provides a good overview :



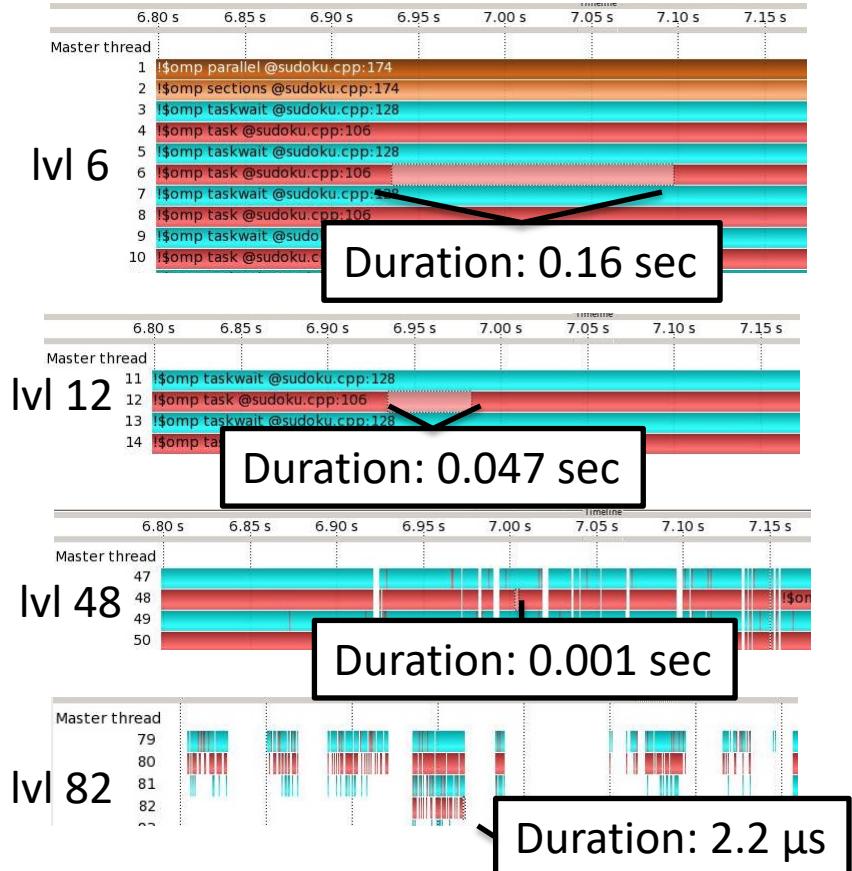
Every thread is executing ~1.3m tasks...



... in ~5.7 seconds.

=> average duration of a task is ~4.4 µs

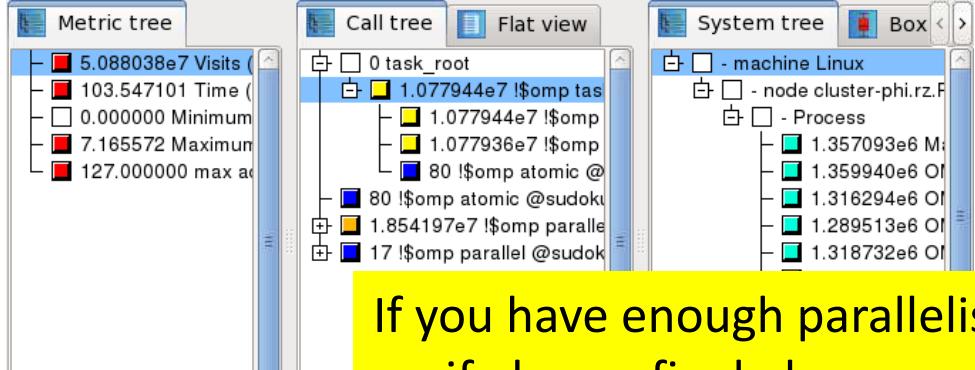
Tracing provides more details:



Tasks get much smaller down the call-stack.

Performance Analysis

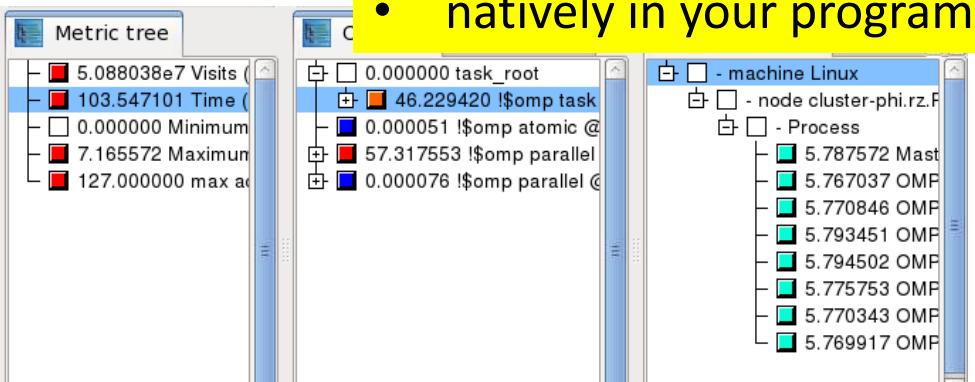
Event-based profiling provides a good overview :



If you have enough parallelism, stop creating more tasks!!

Every thread is doing:

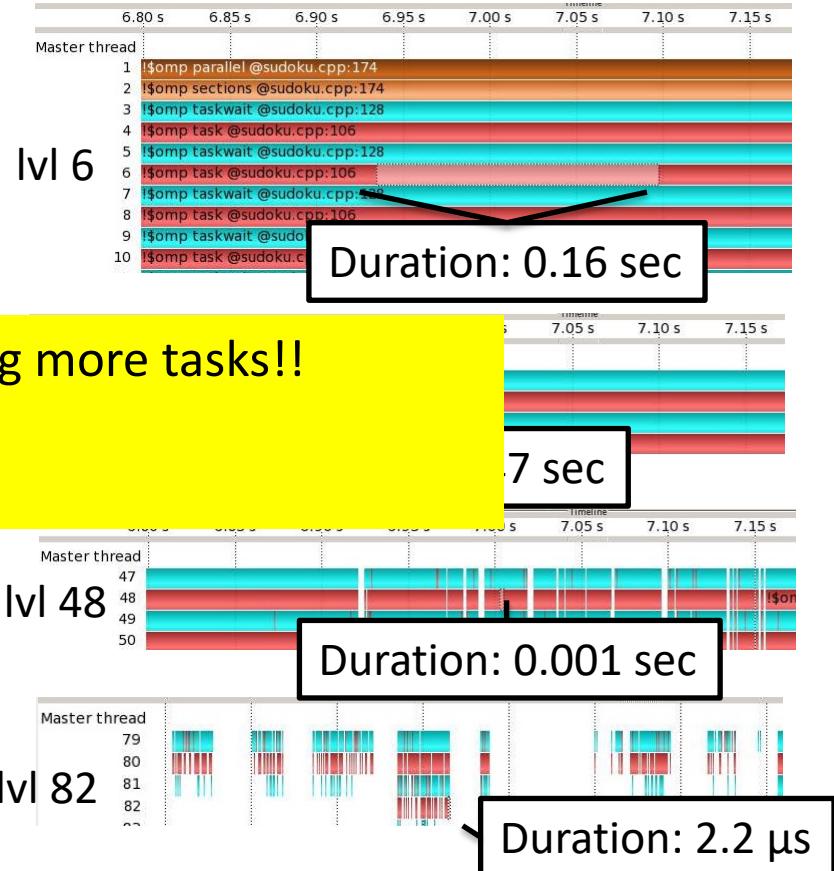
- if-clause, final-clause, mergeable-clause
- natively in your program code



... in ~5.7 seconds.

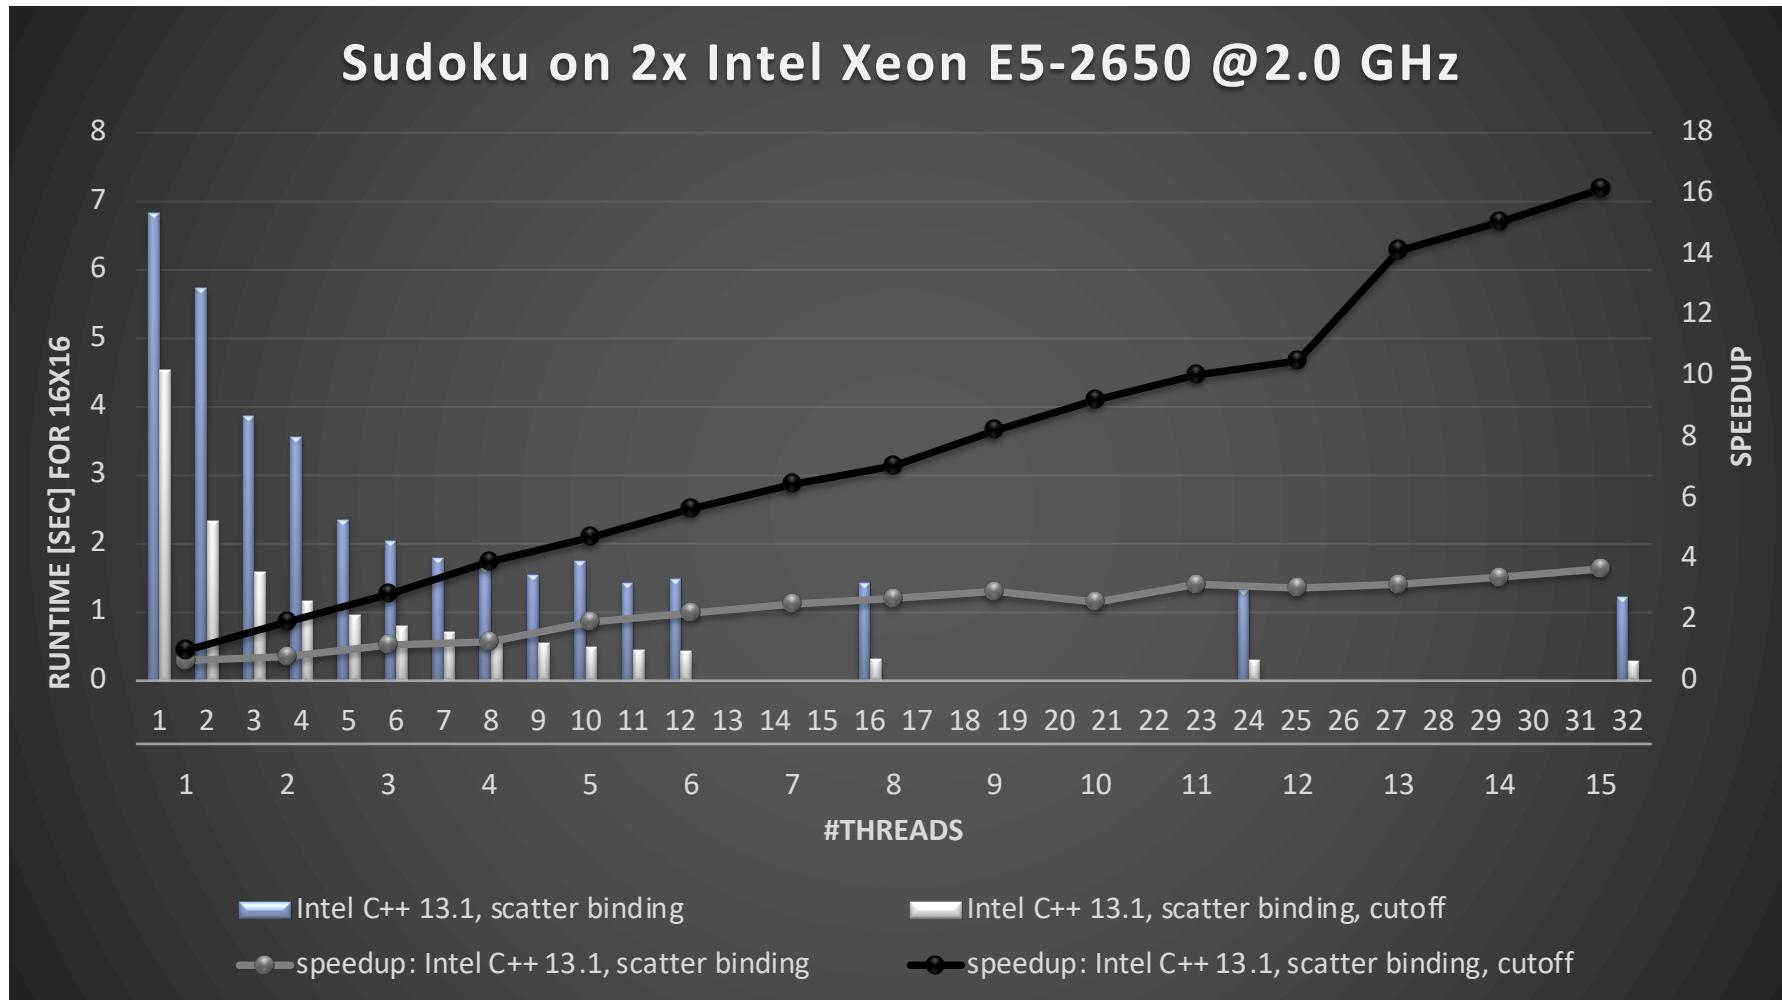
=> average duration of a task is ~4.4 µs

Tracing provides more details:



Tasks get much smaller down the call-stack.

Performance Evaluation (with cutoff)



The if clause

- Rule of thumb: the `if (expression)` clause as a “switch off” mechanism
 - Allows lightweight implementations of task creation and execution but it reduces the parallelism
- If the expression of the `if` clause evaluates to `false`
 - the encountering task is suspended
 - the new task is executed immediately (task dependences are respected!!)
 - the encountering task resumes its execution once the new task is completed
 - This is known as *undelayed task*
- Even if the expression is false, data-sharing clauses are honored

```
int foo(int x) {  
    printf("entering foo function\n");  
    int res = 0;  
    #pragma omp task shared(res) if(false)  
    {  
        res += x;  
    }  
    printf("leaving foo function\n");  
}
```

Really useful to debug tasking applications!



The final clause

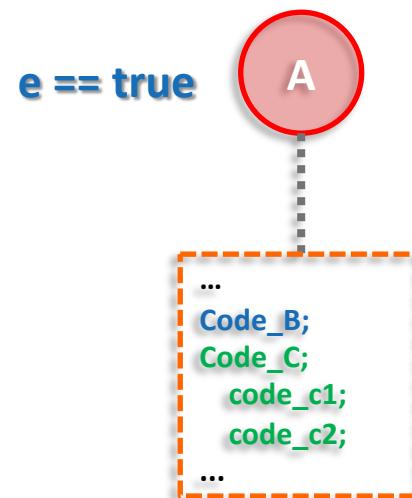
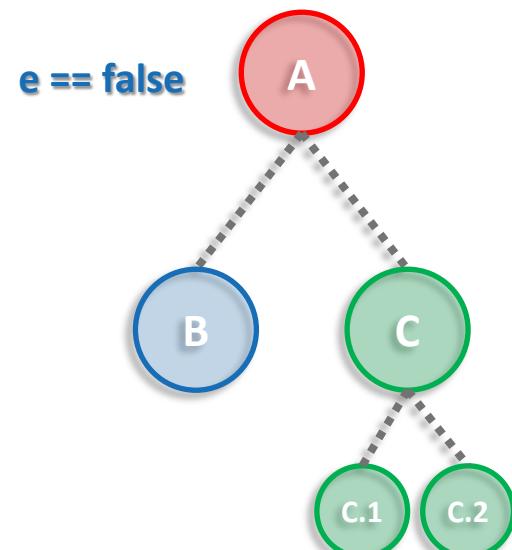
■ The final (expression) clause

- Nested tasks / recursive applications
- allows to avoid future task creation → reduces overhead but also reduces parallelism

■ If the expression of the final clause evaluates to true

- The new task is created and executed normally but in its context all tasks will be executed immediately by the same thread (*included tasks*)

```
#pragma omp task final(e)
{
    #pragma omp task
    { ... }
    #pragma omp task
    { ... #C.1; #C.2 ... }
    #pragma omp taskwait
}
```



■ Data-sharing clauses are honored too!

The mergeable clause

■ The mergeable clause

- Optimization: get rid of “data-sharing clauses are honored”
- This optimization can only be applied in *undeferred* or *included* tasks

■ A Task that is annotated with the mergeable clause is called a *mergeable task*

- A task that may be a *merged task* if it is an *undeferred task* or an *included task*

■ A *merged task* is:

- A task for which the data environment (inclusive of ICVs) may be the same as that of its generating task region

■ A good implementation could execute a merged task without adding any OpenMP-related overhead

Unfortunately, there are no OpenMP commercial implementations taking advantage of `final` neither `mergeable` =()



Programming OpenMP

Using OpenMP Compilers

Christian Terboven

Michael Klemm



Production Compilers w/ OpenMP Support

- GCC
- clang/LLVM
- Intel Classic and Next-gen Compilers
- AOCC, AOMP, ROCmCC
- IBM XL
- ... and many more

- See <https://www.openmp.org/resources/openmp-compilers-tools/> for a list

Compiling OpenMP

- Enable OpenMP via the compiler's command-line switches

- GCC: -fopenmp
- clang: -fopenmp
- Intel: -fopenmp or -qopenmp (classic) or -fiopenmp (next-gen)
- AOCC, AOCL, ROCmCC: -fopenmp
- IBM XL: -qsmp=omp

- Switches have to be passed to both compiler and linker:

```
$ gcc [...] -fopenmp -o matmul.o -c matmul.c
$ gcc [...] -fopenmp -o matmul matmul.o
$./matmul 1024
Sum of matrix (serial): 134217728.000000, wall time 0.413975, speed-up 1.00
Sum of matrix (parallel): 134217728.000000, wall time 0.092162, speed-up 4.49
```

Programming OpenMP

Hands-on Exercises

Christian Terboven

Michael Klemm



Webinar Exercises

- We have implemented a series of small hands-on examples that you can use and play with.
 - Download: `git clone https://github.com/cterboven/OpenMP-tutorial-PRACE-2022.git`
 - Build: `make` (in the corresponding subdirectories)
 - You can then find the compiled code in the “bin” folder to run it
 - We use the GCC compiler mostly, some examples require Intel’s Math Kernel Library
- Each hands-on exercise has a folder “solution”
 - It shows the OpenMP directive that we have added
 - You can use it to cheat ☺, or to check if you came up with the same solution
- Also provided: basic exercises in the `openmp-simple-exercises.tar` archive
 - Instructions contained in the archive: `Exercises_OMP_2021.pdf`

Programming OpenMP

Hands-on Exercises

Christian Terboven

Michael Klemm



Solution of Homework Assignments

Example: sin-cos

```
double do_some_computation(int i) {
    double t = 0.0; int j;
    for (j = 0; j < i*i; j++) {
        t += sin((double)j) * cos((double)j);
    }
    return t;
}

int main(int argc, char* argv[]) {
    const int dimension = 500;
    int i;
    double result = 0.0;
    double t1 = omp_get_wtime();
    #pragma omp parallel for schedule(dynamic) reduction(+:result)
    for (i = 0; i < dimension; i++) {
        result += do_some_computation(i);
    }
    double t2 = omp_get_wtime();
    printf("Computation took %.3lf seconds.\n", t2 - t1);
    printf("Result is %.3lf.\n", result);
    return 0;
}
```

Example: matmul

```
void matmul_seq(double * C, double * A, double * B, size_t n) { ... }

void matmul_par(double * C, double * A, double * B, size_t n) {
#pragma omp parallel for shared(A,B,C) firstprivate(n) \
    schedule(static) // collapse(2)
    for (size_t i = 0; i < n; ++i) {
        for (size_t k = 0; k < n; ++k) {
            for (size_t j = 0; j < n; ++j) {
                C[i * n + j] += A[i * n + k] * B[k * n + j];
            }
        }
    }
}

void init_mat(double * C, double * A, double * B, size_t n) { ... }

void dump_mat(double * mtx, size_t n) { ... }
double sum_mat(double * mtx, size_t n) { ... }

int main(int argc, char *argv[]) { ... }
```

Example: cholesky

```
void cholesky(int ts, int nt, double* Ah[nt][nt]) {  
    for (int k = 0; k < nt; k++) {  
        LAPACKE_dpotrf(LAPACK_COL_MAJOR, 'L', ts, Ah[k][k], ts);  
  
        #pragma omp parallel for  
        for (int i = k + 1; i < nt; i++) {  
            cblas_dtrsm(CblasColMajor, CblasRight, CblasLower, CblasTrans,  
                         CblasNonUnit, ts, ts, 1.0, Ah[k][k], ts, Ah[k][i], ts);  
        }  
  
        #pragma omp parallel for  
        for (int i = k + 1; i < nt; i++) {  
            for (int j = k + 1; j < i; j++) {  
                cblas_dgemm(CblasColMajor, CblasNoTrans, CblasTrans, ts, ts, ts, -1.0,  
                            Ah[k][i], ts, Ah[k][j], ts, 1.0, Ah[j][i], ts);  
            }  
            cblas_dsyrk(CblasColMajor, CblasLower, CblasNoTrans, ts, ts, -1.0,  
                        Ah[k][i], ts, 1.0, Ah[i][i], ts);  
        }  
    }  
}
```

Blocked matrix
w/ block size ts